



Brian Bønk Rueløkke

Principal & Enterprise arkitekt, Data & Analytics

Fellowmind



<https://linkedin.com/in/brianbonk>



<https://brianbonk.dk>



Microsoft

FastTrack Recognized
Solution Architect
Power BI
2022 >>



Microsoft

Certified Trainer
Data Platform

2018 >>

Agenda

Quick warm up demo

What is Kusto / ADX + free stuff

The marriage with Power BI

Demo



QUICK WARM UP

DEMO

Live coding
(hopefully no demo-ghost 👻)

Kusto / ADX / SDX

Azure Data
Explorer

Synapse
Data
Explorer

Fabric
KQL
database

Kusto engine

The Fabric experience for Kusto

Fabric
KQL
database

Proprietary



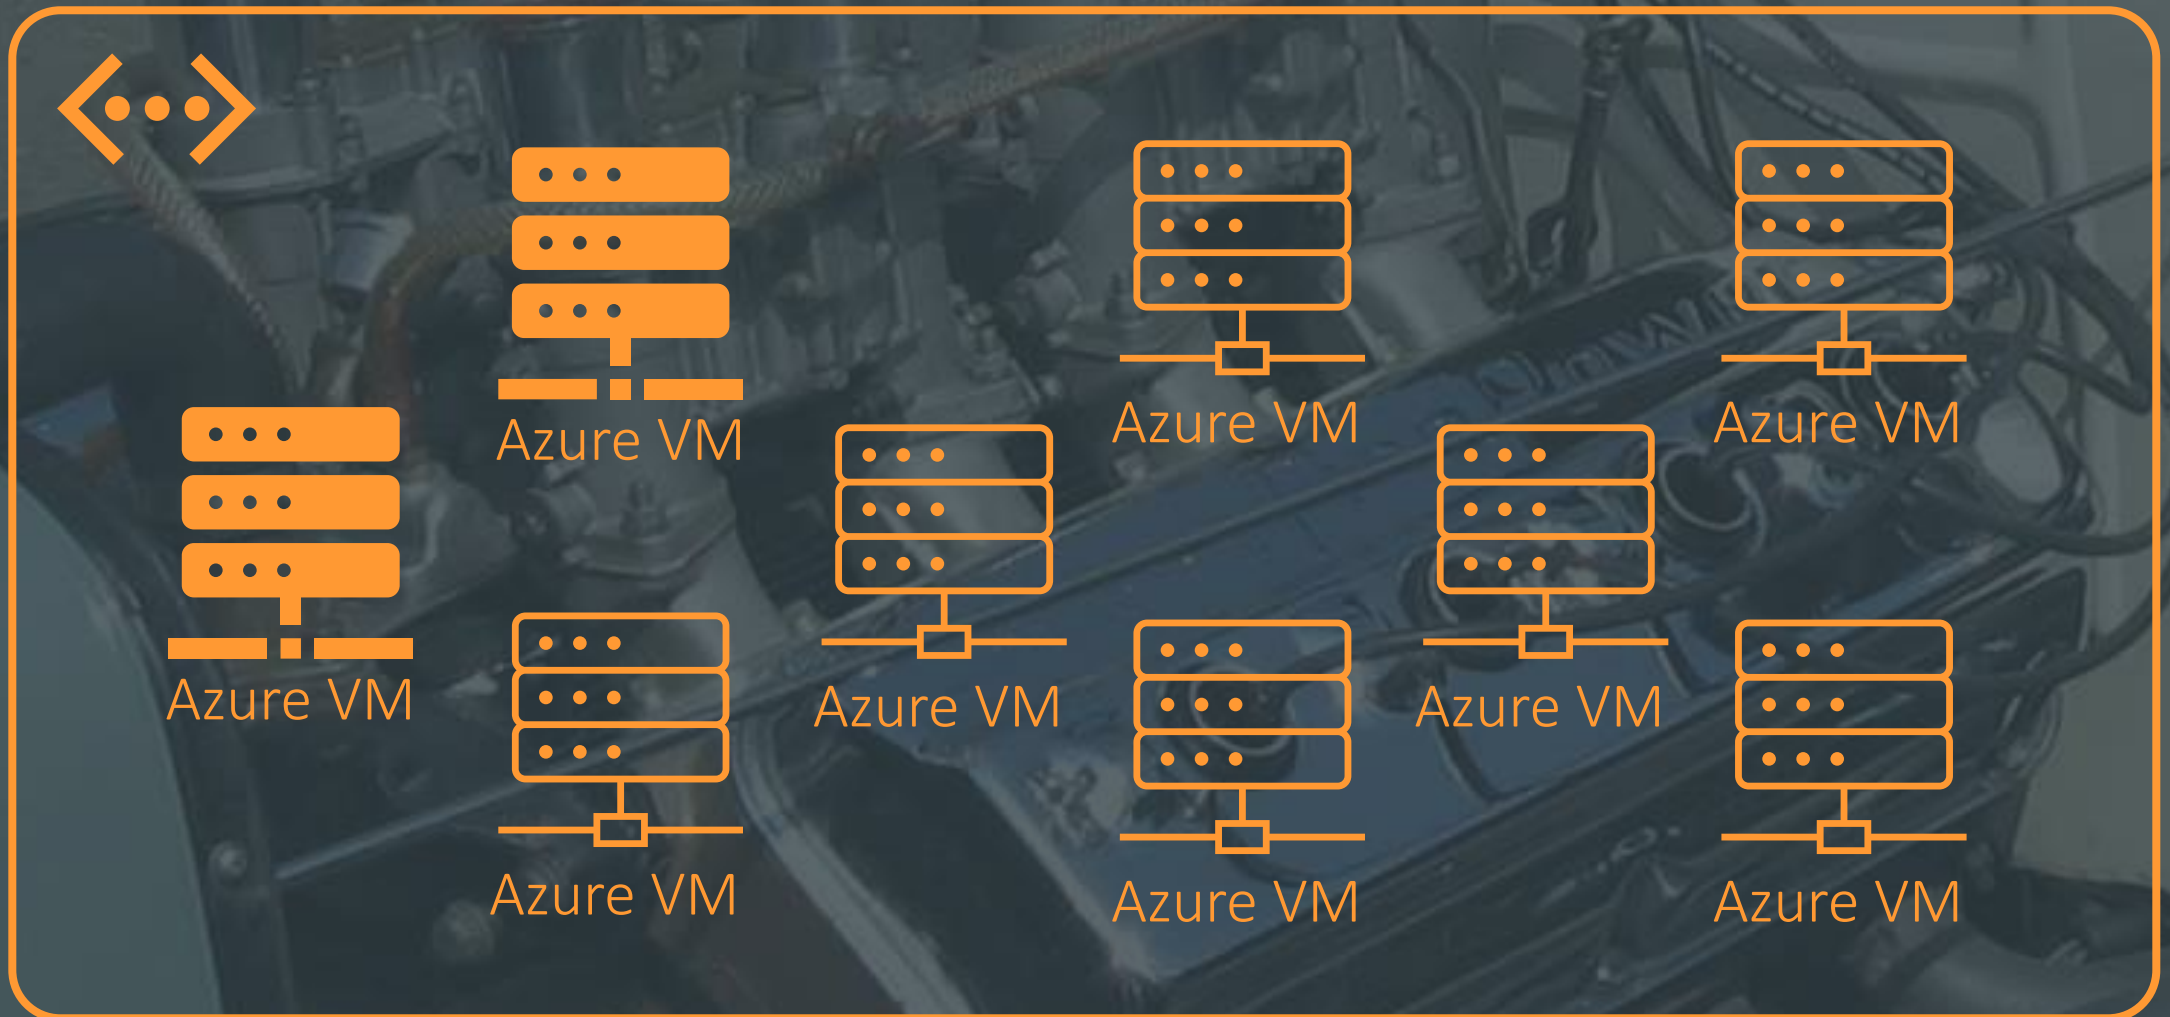
Sync to OneLake



Kusto / ADX / SDX

- Distributed engine on several disks (SSD or spinning)
- Columnar storage → read only needed columns and not all data
- Clusters are Azure VMs → highly parallel
 - For Fabric we don't know yet
- Read-only, delete rarely and no updates
- Fast ingestion → no consistency checks

Kusto / ADX / SDX



Kusto / ADX / SDX



Admin node:
responsible for maintaining the overall cluster metadata

Query Head:
Responsible for accepting and processing Kusto query, when you see, Kusto engine or Kusto query planer, usually refer to query head node.

Data Node:
The most common role, like its name indicates, this node is responsible for: first. storing data; second. contribute the CPU and memory when executing the Kusto query.

Gateway Node:
Responsible for processing external API calls, authentication, and request dispatches.

Fabric architecture for Kusto

Fabric
KQL
database

Proprietary



That's all we know right now



Get started for free

<https://dataexplorer.azure.com/freecluster>

Data modelling Kusto in Power BI

Some standard guidelines of data modelling does not apply

- Single table reporting can be a good option, if you can include all columns from dimensions to the table
- M:M relations are hard to avoid, but not a big deal → all queries will be translated to KQL

Customer
CustomerKey
Name
Birth Date
Marital Status
Education
Occupation
Continent
CountryRegion

Sales
ProductKey
Quantity
Delivery Date
CustomerKey
Sales Amount
Margin
Margin %

Date
Date
Year
Month
MonthNumber
Day of Week
DayNumber

Purchases
Delivery Date
Order Date
ProductKey
Quantity
Invoice Cost
Purchase Amount


Product
ProductKey
Product Name
Price
Color

Data modelling Kusto in Power BI

Star schema is fine if you follow some guidelines:

- All dimensions must be tagged with “IsDimension=true” – demo later
- Use only measures (or columns from dimensions) – do not use columns from fact tables
- Dimensions can be imported if they are <1 mio rows.

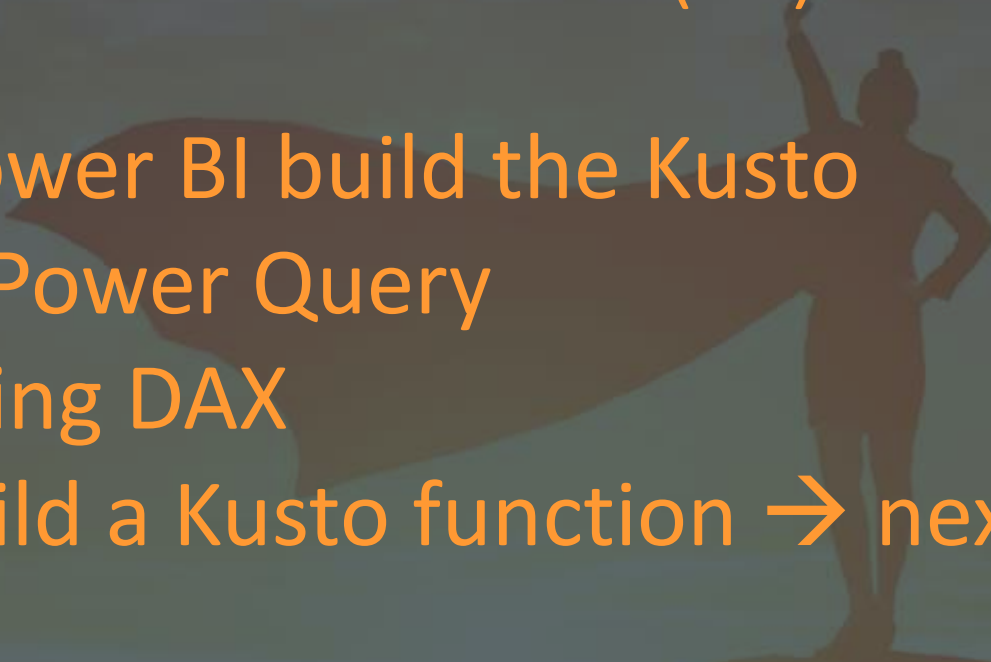
Customer
CustomerKey
Name
Birth Date
Marital Status
Education
Occupation
Continent
CountryRegion



Harness the Power (BI) of Kusto (KQL)

Let Power BI build the Kusto

- In Power Query
- Using DAX
- Build a Kusto function → next slide



Harness the Power (BI) of Kusto (KQL)

```
.create-or-alter function GetSysLogs(TimeWindow:string , Bucket:string )  
{  
cluster('help').database('SampleLogs').RawSysLogs  
| where timestamp > ago(totimespan(TimeWindow))  
| summarize LogCount=count() by name, bin(timestamp, totimespan(Bucket))  
| order by timestamp asc  
}
```

```
GetSysLogs('5d','1h')
```

Power BI Builds the query

From Power Query From Filters From Relations From Visuals

```
[ "Customers" ]
| where [ "RegionCountryName" ] == "United States"
| project-rename [ "semijoin1.c5" ] = [ "CustomerKey" ], [ "semijoin1.c27" ] =
[ "StateProvinceName" ]
| join hint.strategy=broadcast kind=inner ( [ "SalesFact" ]
| project [ "SalesAmount" ], [ "TotalCost" ], [ "CustomerKey" ]
| extend [ "Margin" ] = ( [ "SalesAmount" ] ) -
( [ "TotalCost" ] ), [ "t0_0" ] = tolong( [ "SalesAmount" ] )
| project-rename [ "basetable0.c5" ] = [ "CustomerKey" ], [ "basetable0.a0" ] = [ "t0_0" ],
[ "basetable0.a1" ] = [ "t2_0" ]) on $left.[ "semijoin1.c5" ] == $right.[ "basetable0.c5" ]
| summarize [ "a0" ] = sum( [ "basetable0.a0" ] ), [ "a1" ] = sum( [ "basetable0.a1" ] ) by
[ "semijoin1.c27" ]
```

The marriage with Power BI

DEMO

Live coding
(hopefully no demo-ghost 👻)



Monitor your Kusto queries

- Use `.show queries` → demo later
- Separate queries by user
- Separate queries by report
- Pay attention to CPU consumption and not just duration

Monitor your queries



DEMO

Live coding
(hopefully no demo-ghost 👻)

.show queries

```
.show commands-and-queries
// | where ClientActivityId startswith "KPBI"
| where StartedOn > datetime(2022-12-04T09:04:51.198128Z)
| where User ==current_principal_details().UserPrincipalName
| order by StartedOn asc
| extend delay=datetime_diff("Millisecond",next(StartedOn),LastUpdatedOn)
| extend delay=iff(delay<0 or delay> 5000,0,delay)
| extend MB=format_bytes(MemoryPeak)
| extend Isgetschema=Text has "getschema"
| extend IsPreview=Text has "limit 1000 "
| extend Len=strlen(Text)
| extend TextLength=strlen(Text)
| extend ScannedData=format_bytes(tolong(CacheStatistics.Shards.Hot.HitBytes))
| fork
    Queries=(where CommandType == "Query" and Isgetschema==false and IsPreview ==false | project StartedOn, LastUpdatedOn, Duration, TotalCpu,
ScannedData, MB,Text | order by StartedOn asc)
    Commands=(where CommandType == "AdminThenQuery" | project StartedOn,LastUpdatedOn, Duration, TotalCpu, MB, ScannedData ,Text | order by StartedOn
asc)
    Detail=(project StartedOn,State,FailureReason, Duration, delay, TotalCpu,Isgetschema, MB, ScannedData,ClientActivityId, Text | order by StartedOn
asc)
    Slow=(where CommandType == "Query" | project Duration, TotalCpu, MB, ScannedData,Text | order by Duration)
    Getschema=(where CommandType == "Query" and Isgetschema==true | project StartedOn,Text)
    Summary=(summarize Commands=countif(CommandType == "AdminThenQuery"),Queries=countif(CommandType=="Query"),DelayCommand=sumif(delay, CommandType
=="AdminThenQuery"),
        DelayQuery=sumif(delay, CommandType == "Query"),mn=min(StartedOn),mx=max( LastUpdatedOn),TotCPU=sum(TotalCpu),TotDuration=sum(Duration)
        | extend OverallDuration=(mx-mn))
```

Thank you!

Brian Bønck Rueløkke

Principal & Enterprise arkitekt, Data & Analytics


Fellowmind



<https://linkedin.com/in/brianbonk>



<https://brianbonk.dk>



Microsoft

FastTrack Recognized
Solution Architect
Power BI
2022 >>



Microsoft

Certified Trainer
Data Platform

2018 >>