

Discover the secrets to supercharging
your analytics with Microsoft Fabric
Warehouse






Brian Bønk Rueløkke

Principal & Enterprise architect, Data & AI

Fellowmind

 <https://linkedin.com/in/brianbonk>
 <https://brianbonk.dk>
 <https://github.com/brianbonk>





**AN OVERVIEW
OF FABRIC
AFTER BUILD**



Intelligent data foundation



Data
Factory



Data
Engineering



Data
Warehouse



Data
Science



Real-Time
Intelligence



Power
BI



Industry
Solutions



Powered by AI with Copilot in Microsoft Fabric



Catalog for data in motion

Real-Time Hub



Unified data foundation

OneLake



Intelligent data foundation



Data
Factory



Data
Engineering



Data
Warehouse



Data
Science



Real-Time
Intelligence



Power
BI



Industry
Solutions



Powered by AI with Copilot in Microsoft Fabric

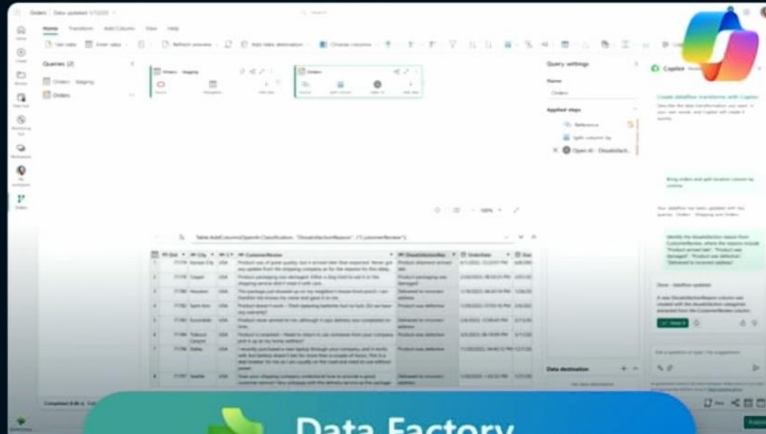


Catalog for data in motion
Real-Time Hub

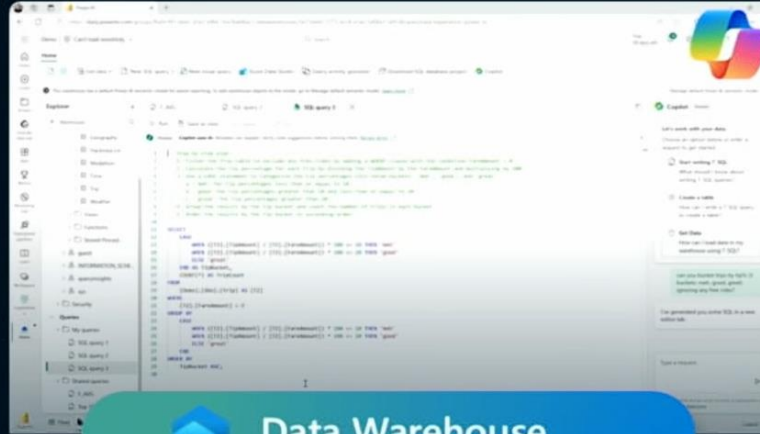


Unified data foundation
OneLake

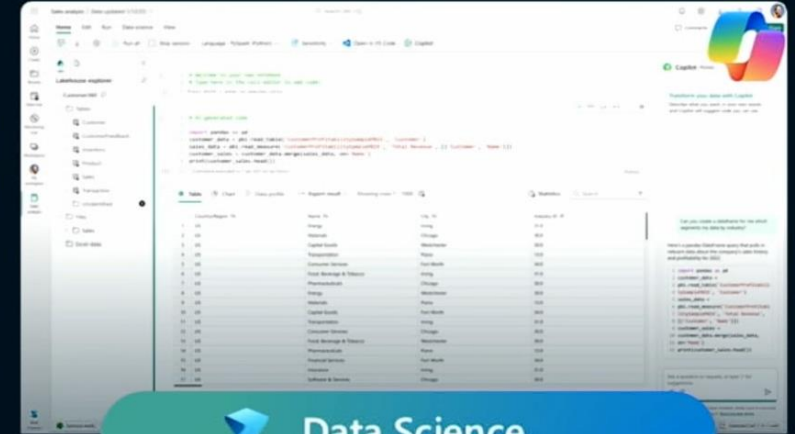
Copilot Integrated in every Microsoft Fabric Experience



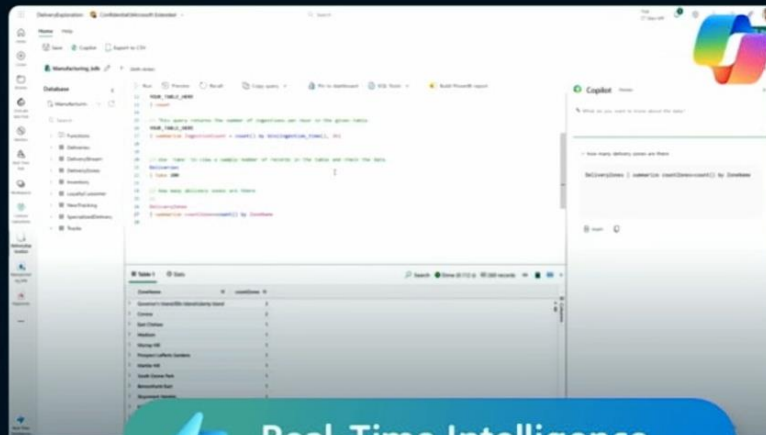
Data Factory



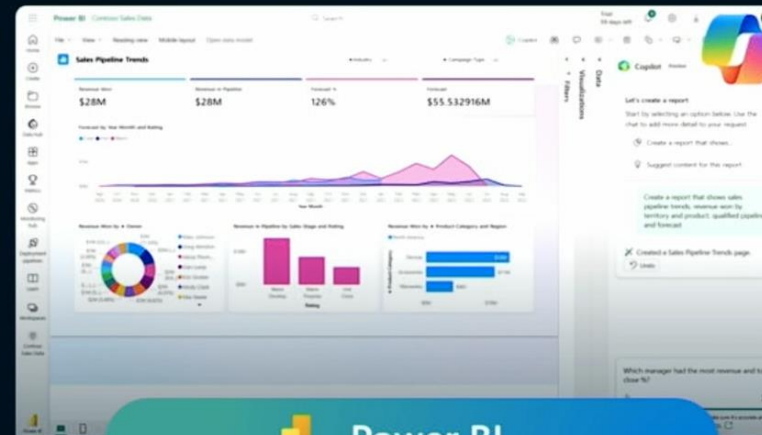
Data Warehouse



Data Science



Real-Time Intelligence



Power BI

API for GraphQL

⋮

Home

+

📁

🔄

👤

👥

🏠

🏢

ContosoOutdoorsAPI

OneLake data hub

Workspaces

Contoso Outdoors

ContosoOutdoorsAPI

ContosoPipeline

PartnerInventory

ContosoSales

Search

11

⚙️

⬇️

?

👤

Home

⚙️

📅 Get data

🔗 Copy endpoint

Schema explorer

Search

productCategories

products

addresses

vProductAndDescr...

productDescriptions

productModelPro...

vProductModelCat...

salesOrderHeaders

customerAddresses

productModels

customers

topsellingproducts

top5Products

Query1

×

+

▶ Run

```
1 query{
2   customers (first: 10) {
3     items{
4       CustomerID,
5       FirstName,
6       LastName
7     }
8   }
9 }
10
```

Query variables

1	
2	{
3	
4	}
5	

Results

```
26 "CustomerID": 5,
27 "FirstName": "Lucy",
28 "LastName": "Harrington"
29 },
30 {
31   "CustomerID": 6,
32   "FirstName": "Rosmarie",
33   "LastName": "Carroll"
```



API for GraphQL

File Edit Selection View Go Run Terminal Help

BuildDemo

EXPLORER

- BUILDEMO
 - .github
 - AppCreationScripts
 - node_modules
 - public
 - src
 - components
 - ChartStyles.css
 - FunctionData.jsx
 - FunctionDataHelloFabric.jsx
 - JS GraphQLAPIData.js
 - GraphQLBarChartComponent.jsx
 - GraphQLData.jsx
 - GraphQLLineChartComponent.jsx
 - GraphQLPieChartComponent.jsx
 - GraphQLTableGridComponent.jsx
 - PageLayout.jsx
 - PieStyles.css
 - ProfileData.jsx
 - SignInButton.jsx
 - SignOutButton.jsx
 - styles
 - App.jsx
 - AppStyles.css
 - authConfig.js
 - dataFunction.js
 - dataFunctionHelloFabric.js
 - fabricEndpointsConfig.js
 - graph.js
 - graphQL.js
 - index.js
 - .gitignore
 - CHANGELOG.md
 - CONTRIBUTING.md
 - LICENSE.md

GraphQLPieChartComponent.jsx M GraphQLBarChartComponent.jsx M GraphQLLineChartComponent.jsx M App.jsx M JS graphQL.js M JS graphQL.js (Working Tree) M PageLayout.jsx

```
src > JS graphQL.js > callGraphQLAPI > options
9 export const graphqlQuery = `query {
15   emailAddress
16   salesOrderHeader {
17     items {
18       SalesOrderID
19       OrderDate
20       TotalDue
21     }
22   }
23 }`
24
25 `;
26
27 export async function callGraphQLAPI(accessToken, query=null, variables=null) {
28   const headers = new Headers();
29   if(query == null) {
30     query = graphqlQuery;
31   }
32   const bearer = `Bearer ${accessToken}`;
33   const graphQLEndpoint = "https://dxtapi.fabric.microsoft.com/v1/workspaces/f4ca5010-490f-433d-b9f9-1738d73bb961/graphqlapis/22aad2df-7862-4aa6-8c67-dde7119489a7/graphq";
34   let data;
35   if(variables == null) {
36     data = JSON.stringify({
37       query: query
38     });
39   } else {
40     data = JSON.stringify({
41       query: query,
42       variables: variables
43     });
44   }
45
46   headers.append("Authorization", bearer);
47   headers.append("Content-Type", 'application/json');
48
49   const options = {
50     method: "POST",
51     headers: headers,
52     body: data
```






API for GraphQL

Contoso Outdoors

Sign Out

Welcome to Contoso Outdoors Portal



GraphQL Customer Orders Details

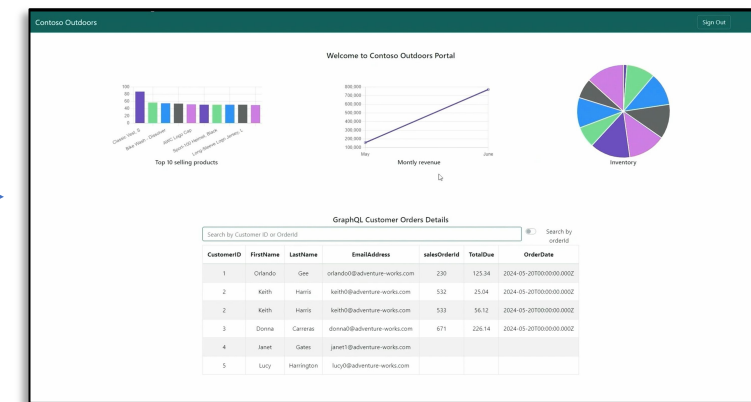
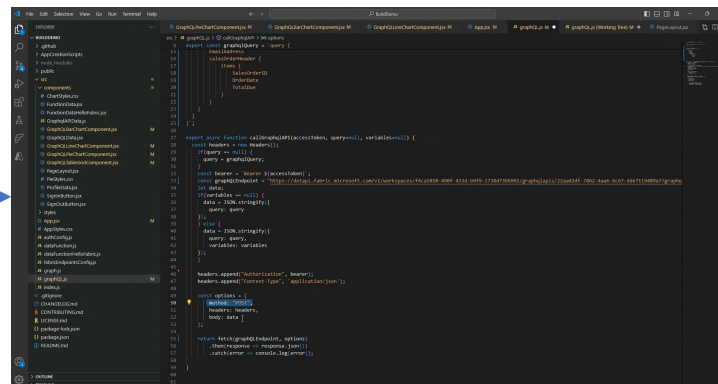
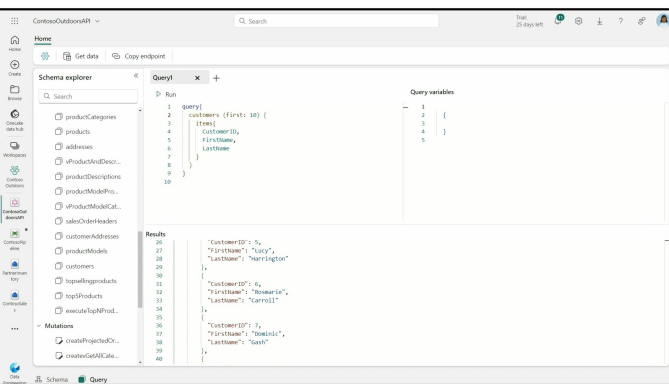
Search by Customer ID or OrderId



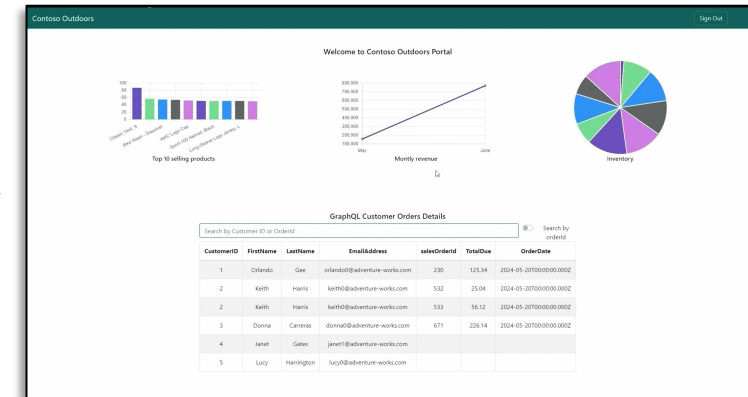
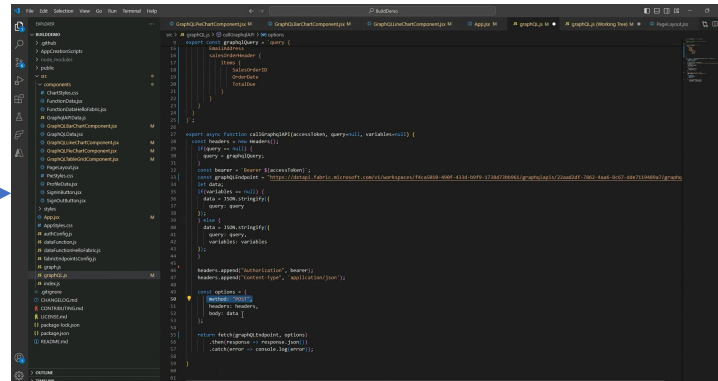
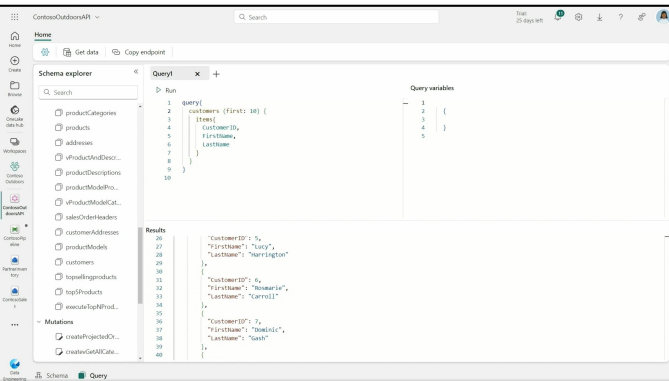
Search by
orderId

CustomerID	FirstName	LastName	EmailAddress	salesOrderId	TotalDue	OrderDate
1	Orlando	Gee	orlando0@adventure-works.com	230	125.34	2024-05-20T00:00:00.000Z
2	Keith	Harris	keith0@adventure-works.com	532	25.04	2024-05-20T00:00:00.000Z

API for GraphQL



API for GraphQL



Takeaways

API for GraphQL in Fabric

Bridging the gap between data and applications



VS



VS



VS



Synapse Data Warehouse

Synapse Serverless SQL

Synapse Dedicated pool

Databricks



VS



Synapse
Serverless SQL

VS



Synapse
Dedicated pool

VS



Databricks

Storage: Fabric focuses on a single storage solution, the Unified Data Location, built on Azure Data Lake Storage Gen2. This eliminates the need for role-based access control and simplifies data management for organizations.

Integration: Seamless integration with multiple Azure services, including Azure Data Factory, Azure Synapse Analytics, Synapse Real-Time Analytics, Synapse Data Warehousing, Power BI and Data Explorer, with automatic provisioning of the underlying hardware.

Collaboration: Dedicated workspaces enable diverse developers – including data engineers and data scientists – to collaborate effortlessly.



Synapse Data
Warehouse

VS



Synapse
Serverless SQL

VS



Synapse
Dedicated pool

VS



Databricks

A unified workspace: Synapse Analytics provides a unified workspace where data engineers and data scientists can collaborate on big data and SQL-based analytics tasks.

Real-time analysis: The integration of Apache Spark and dedicated SQL pool enables users to perform real-time analysis - on both structured and unstructured data.

Security and Governance: Synapse Analytics provides robust security features and fine-grained access controls to ensure data is protected.



Synapse Data
Warehouse

VS



Synapse
Serverless SQL

VS



Synapse
Dedicated pool

VS



Databricks

Scalability: Databricks can easily handle large data tasks as it has the ability to scale horizontally.

Notebooks: Interactive notebooks allow users to execute and visualize code in real time, facilitating data exploration and model training.

Integration: Seamless integration with other Azure services such as Azure Storage and Azure Data Lake Storage simplifies the process of data import and extraction.



Data Warehouse

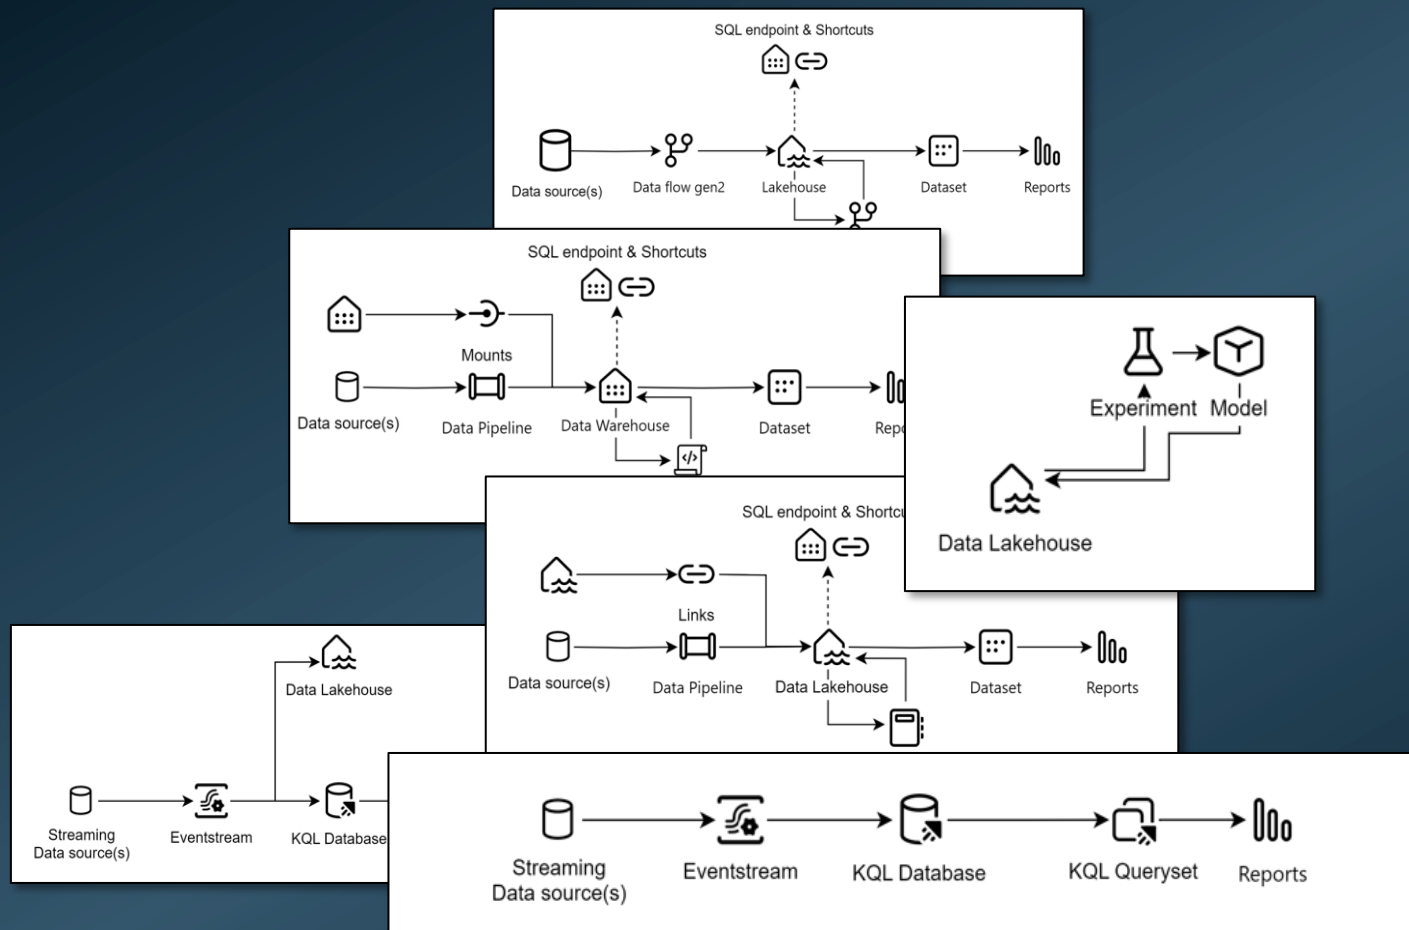


Lakehouse

Warehouse	Lakehouse
Structured	Structured, Semi structured, Unstructured
SQL language	Notebooks
Schemas, Tables	Folders, Files and tables
OLS, RLS, CLS	RLS, TLS
T-SQL	Spark (Scala, PySpark, SparkSQL)



Data
Warehouse



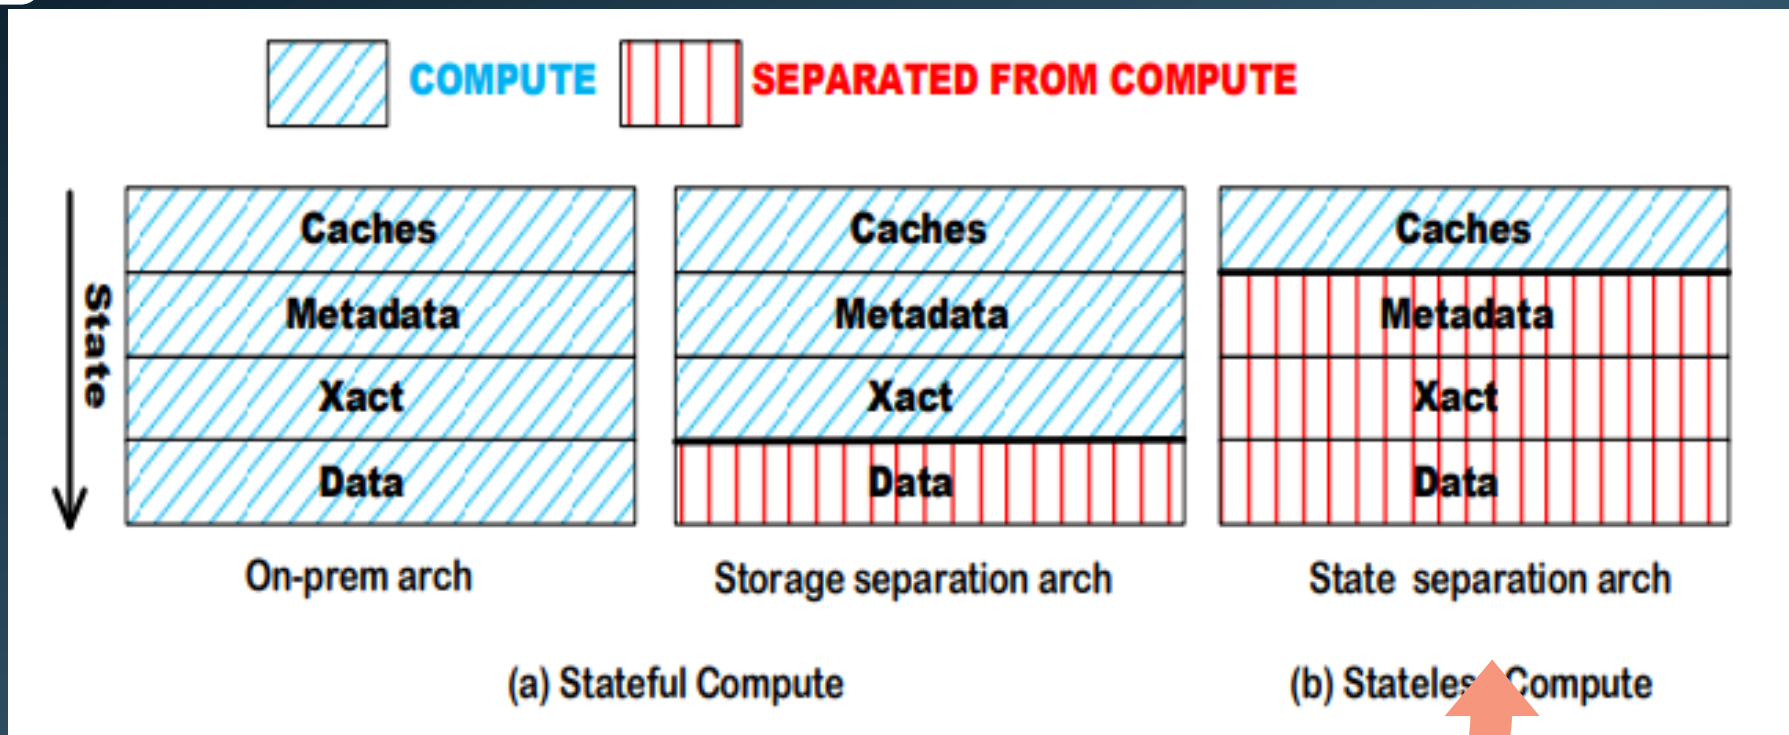


Data
Warehouse



The Polaris engine

Stateless compute



Synapse and Fabric

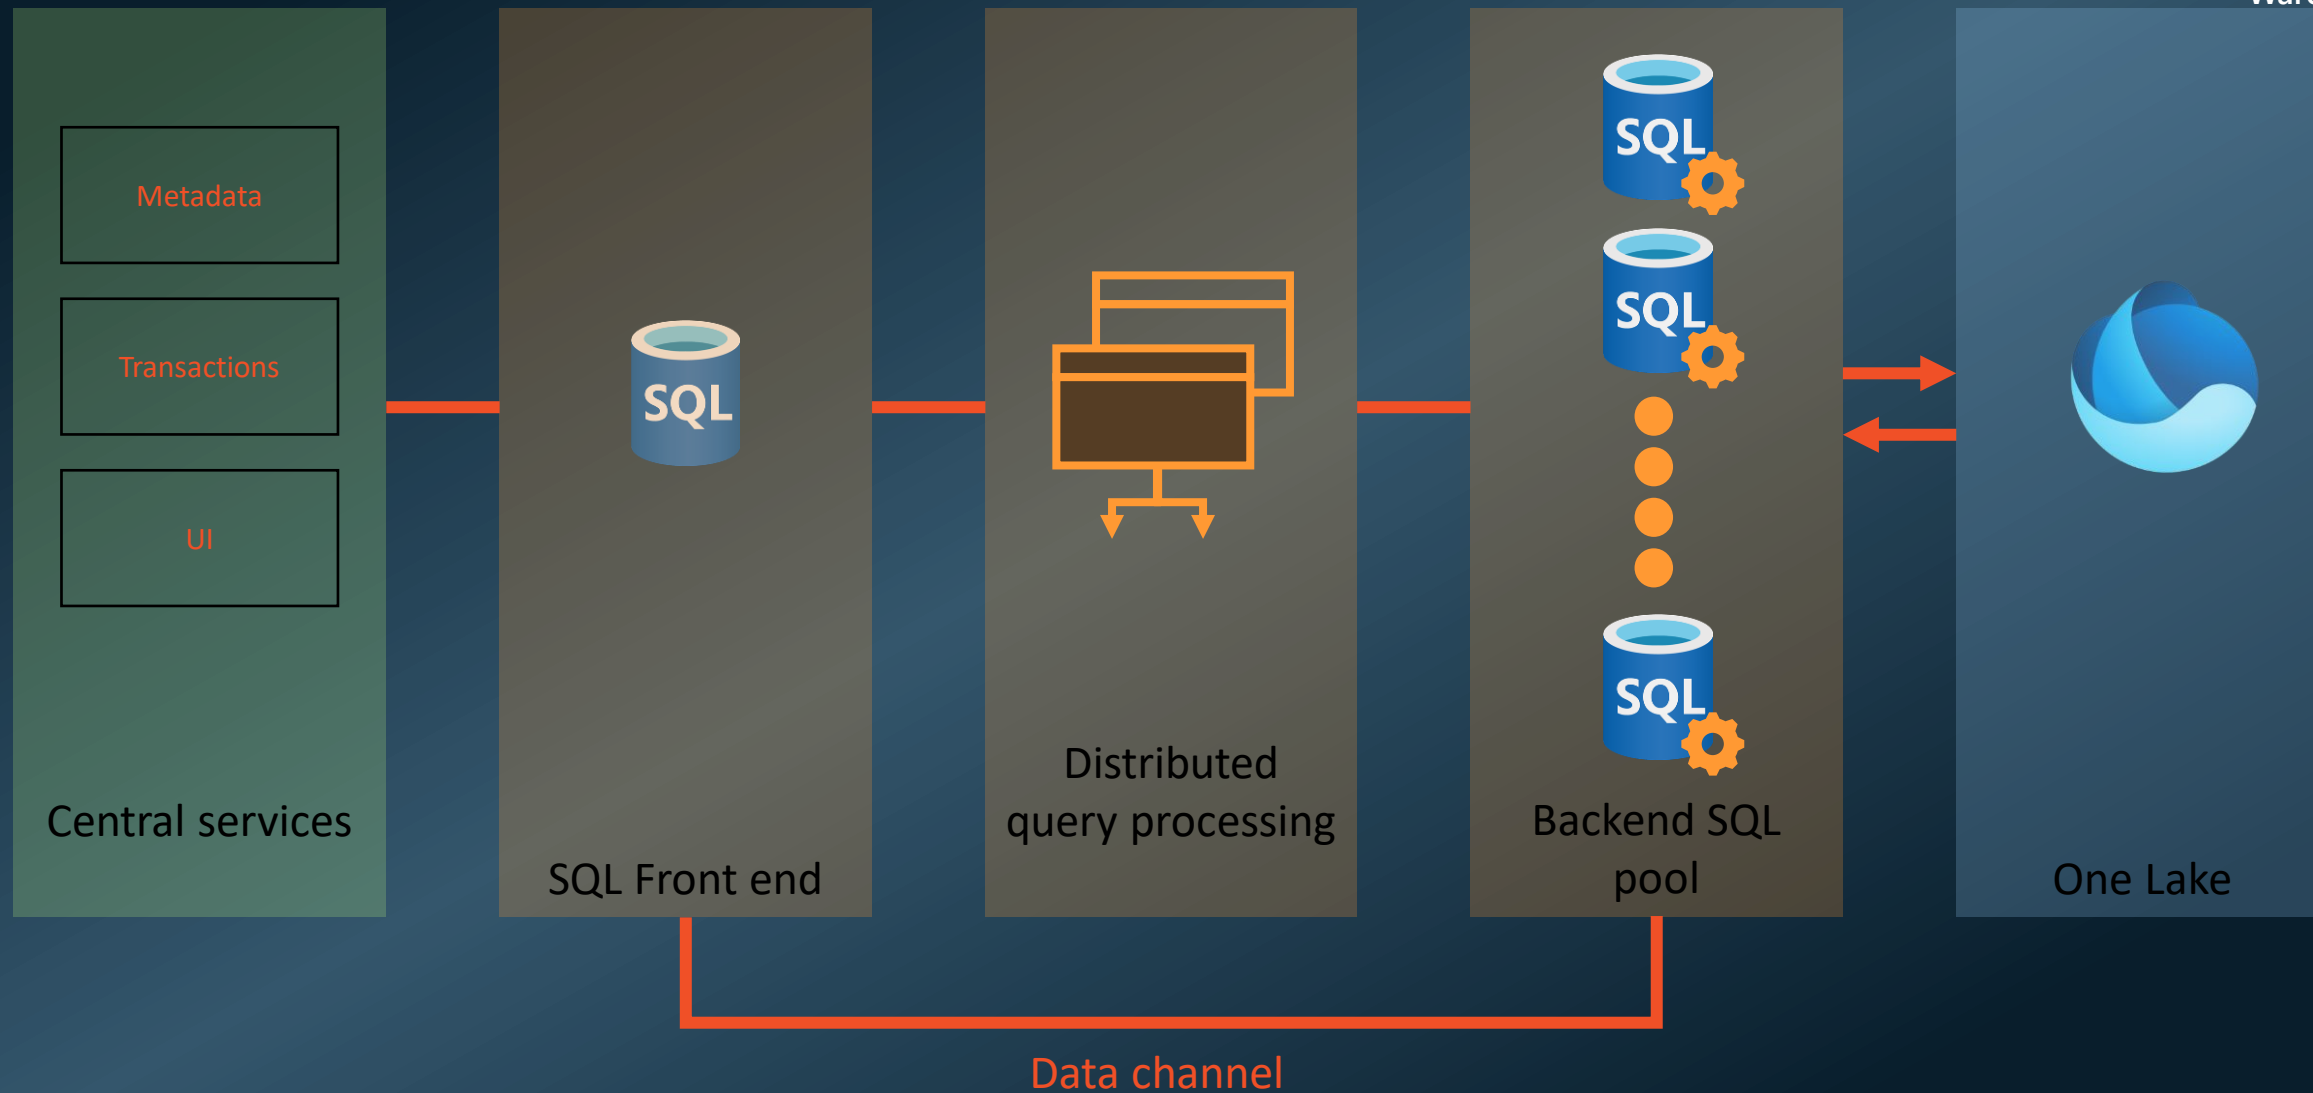


The Polaris engine

Service architecture



Data
Warehouse






T-SQL Surface Area



The T-SQL Surface area covers the supported syntax within both the Warehouse & the Lakehouse SQL Endpoint


Currently unsupported syntax (and the potential impact...)



ALTER TABLE...
...<ADD/ALTER/DROP>

Identity Columns


TRUNCATE



CREATE ROLE/USER

MERGE

Recursive Queries



BULK LOAD

SP_SHOWSPACEUSED

SET TRANSACTION
ISOLATION LEVEL

Transactions & Isolation Levels

Transactions are included (including multi-table)

Only **Snapshot Isolation** supported (Optimistic Concurrency)

Locks are at the **table level**

Use **sys.dm_tran_locks** to show current locks

Schema Stability
Sch-S

SELECT


Intent Exclusive
IX

INSERT
DELETE
UPDATE
COPY INTO

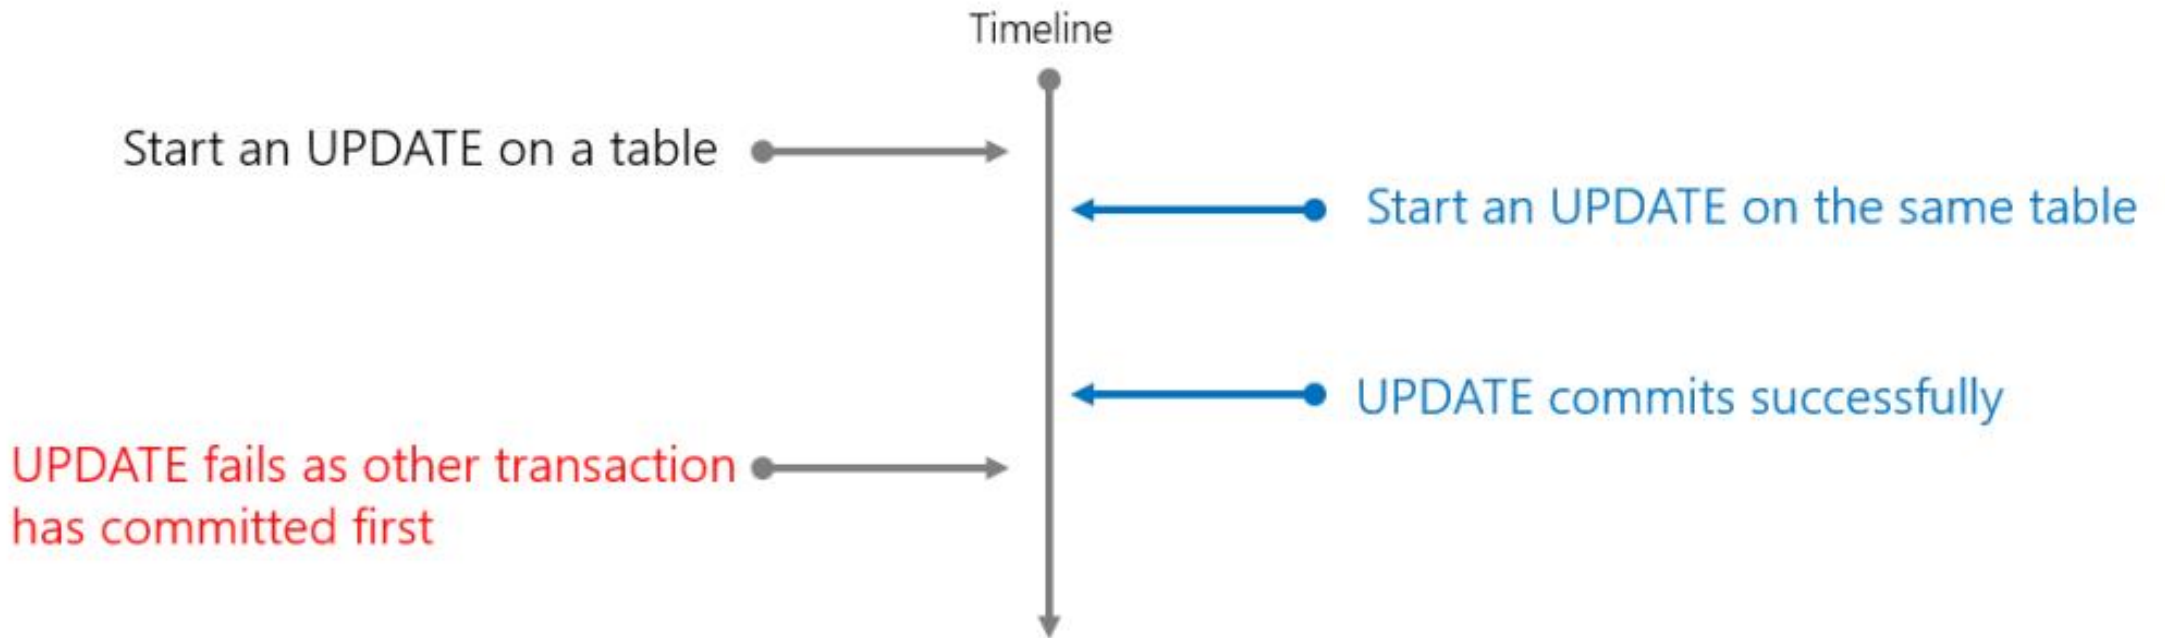
Schema
Modification
Sch-M

DDL

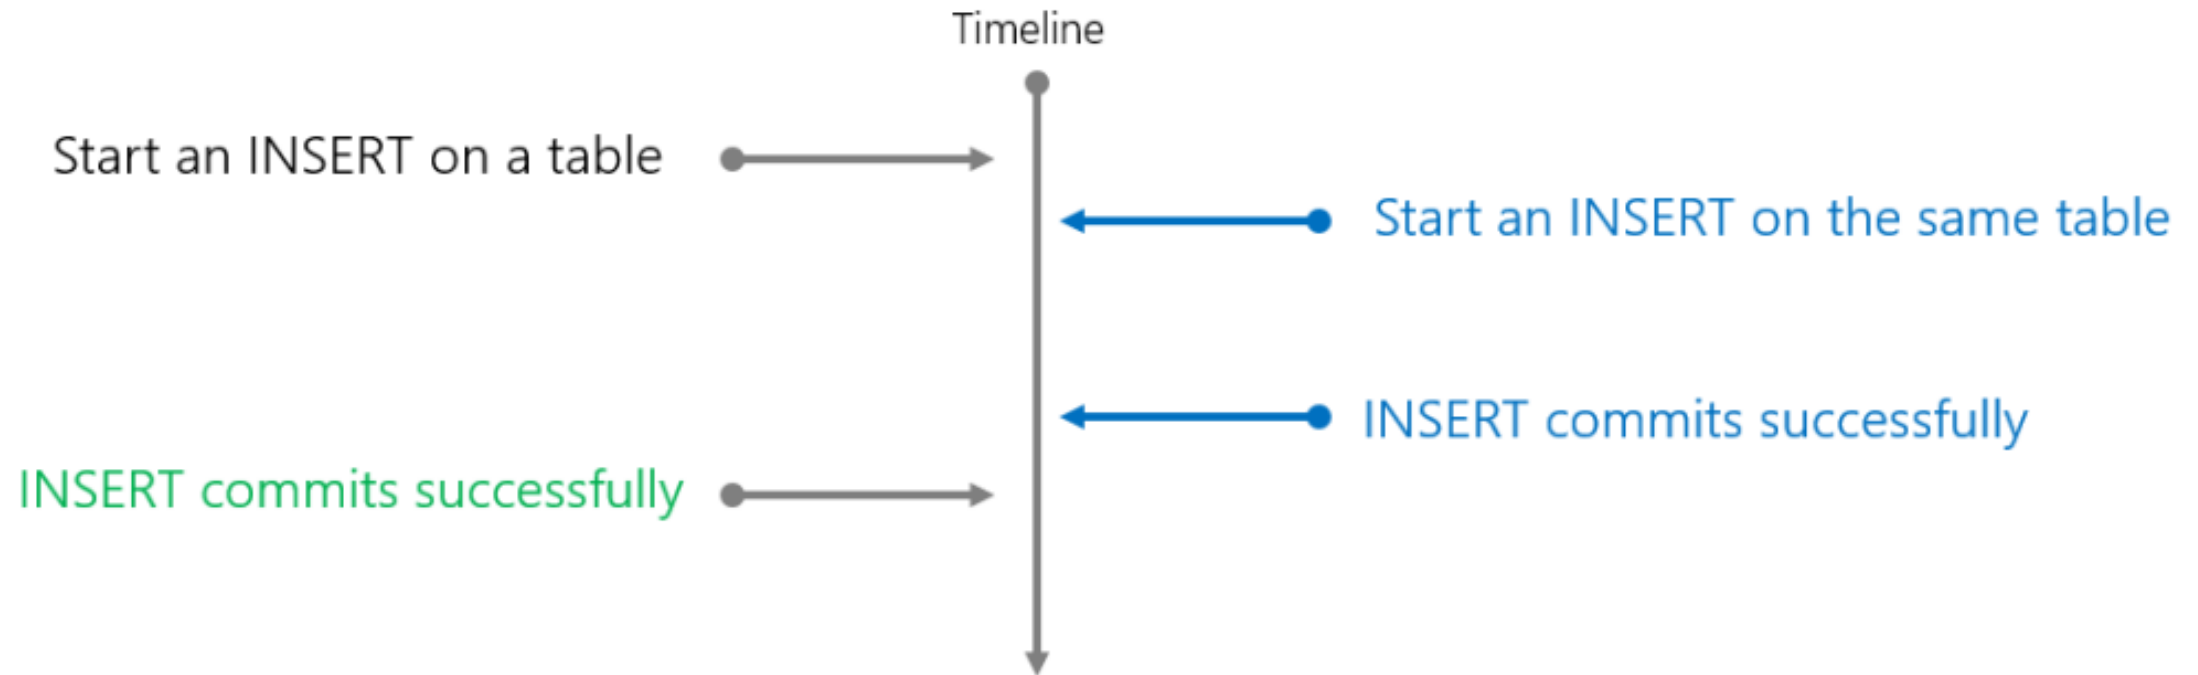
Transactions & Isolation Levels



Transactions & Isolation Levels



Transactions & Isolation Levels





Ingestion: Loading the Warehouse



Data
Warehouse

Code

SQL COPY INTO... <from Azure storage>

SQL CREATE TABLE AS...SELECT

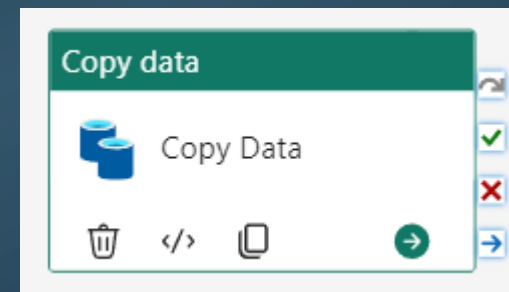
SQL INSERT INTO...VALUES / SELECT

Pyodbc/JDBC (from Notebooks)


Low/No-Code

Pipelines

“Brute force”



Dataflows Gen2
(Power Query)



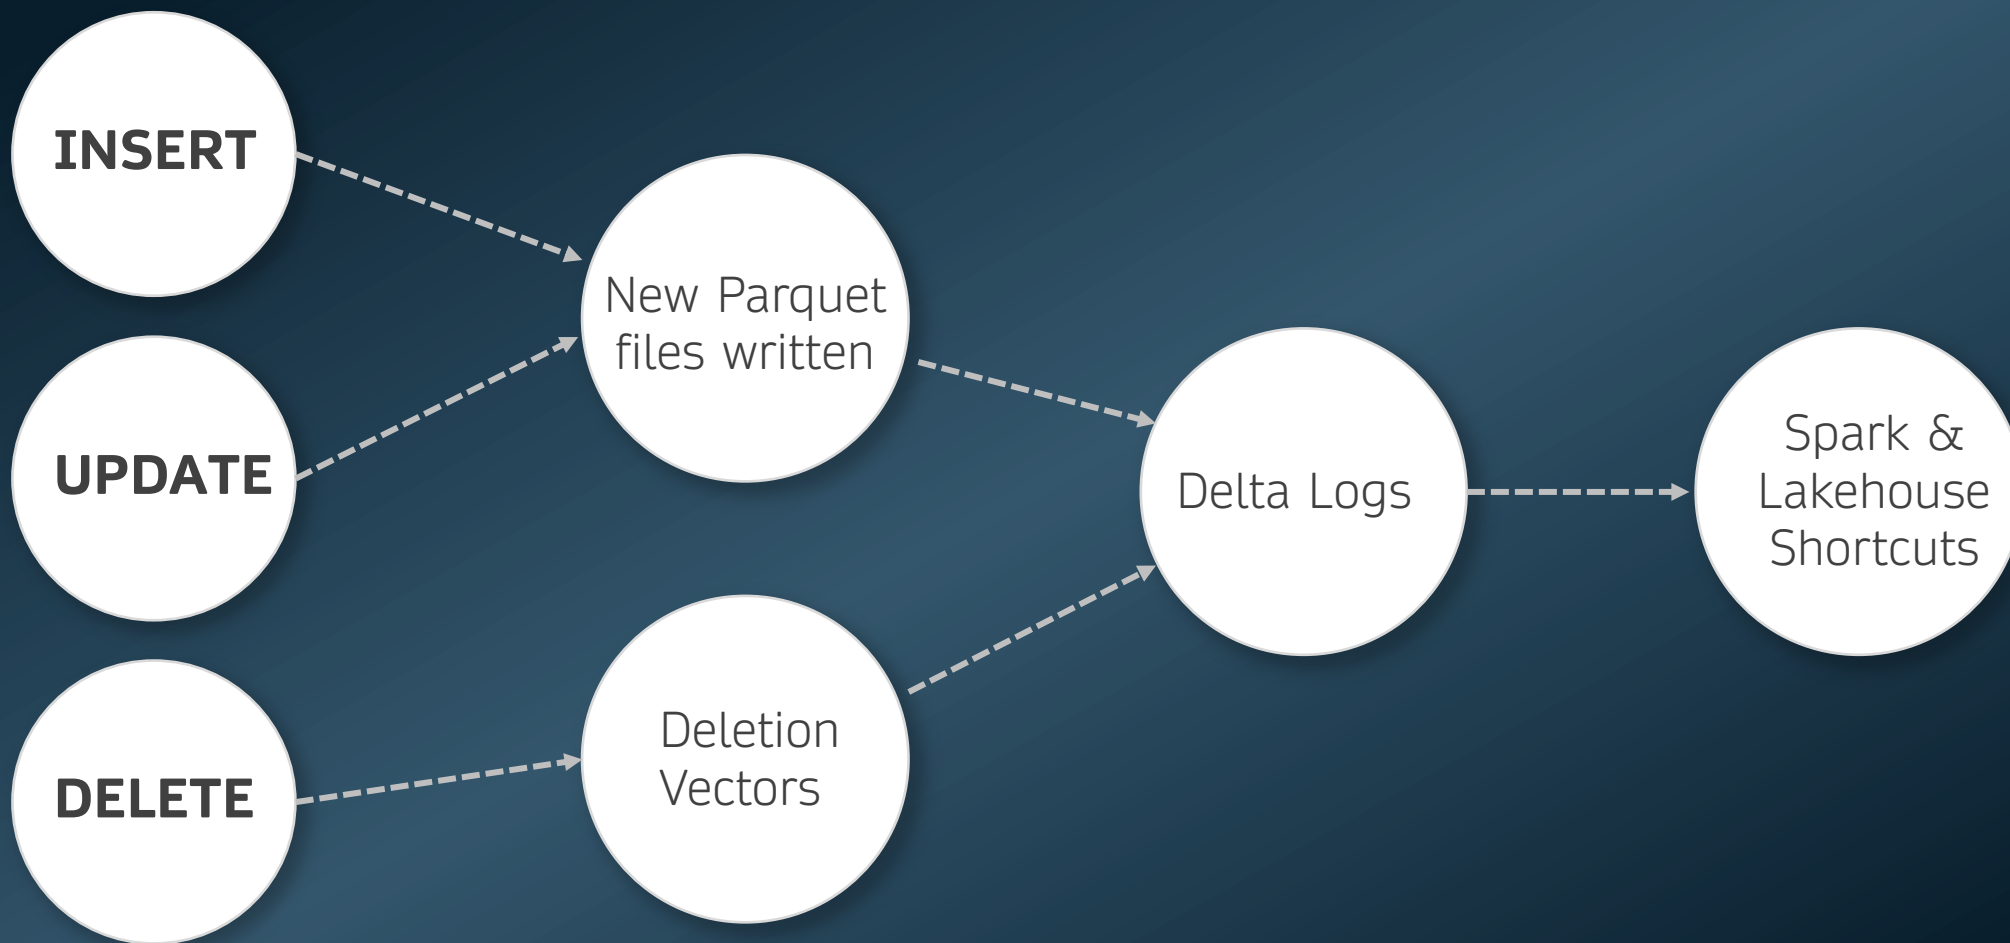


Ingestion: Loading the Warehouse



Data
Warehouse

All SQL operations are performed on Parquet files in OneLake





Ingestion: Loading the Warehouse

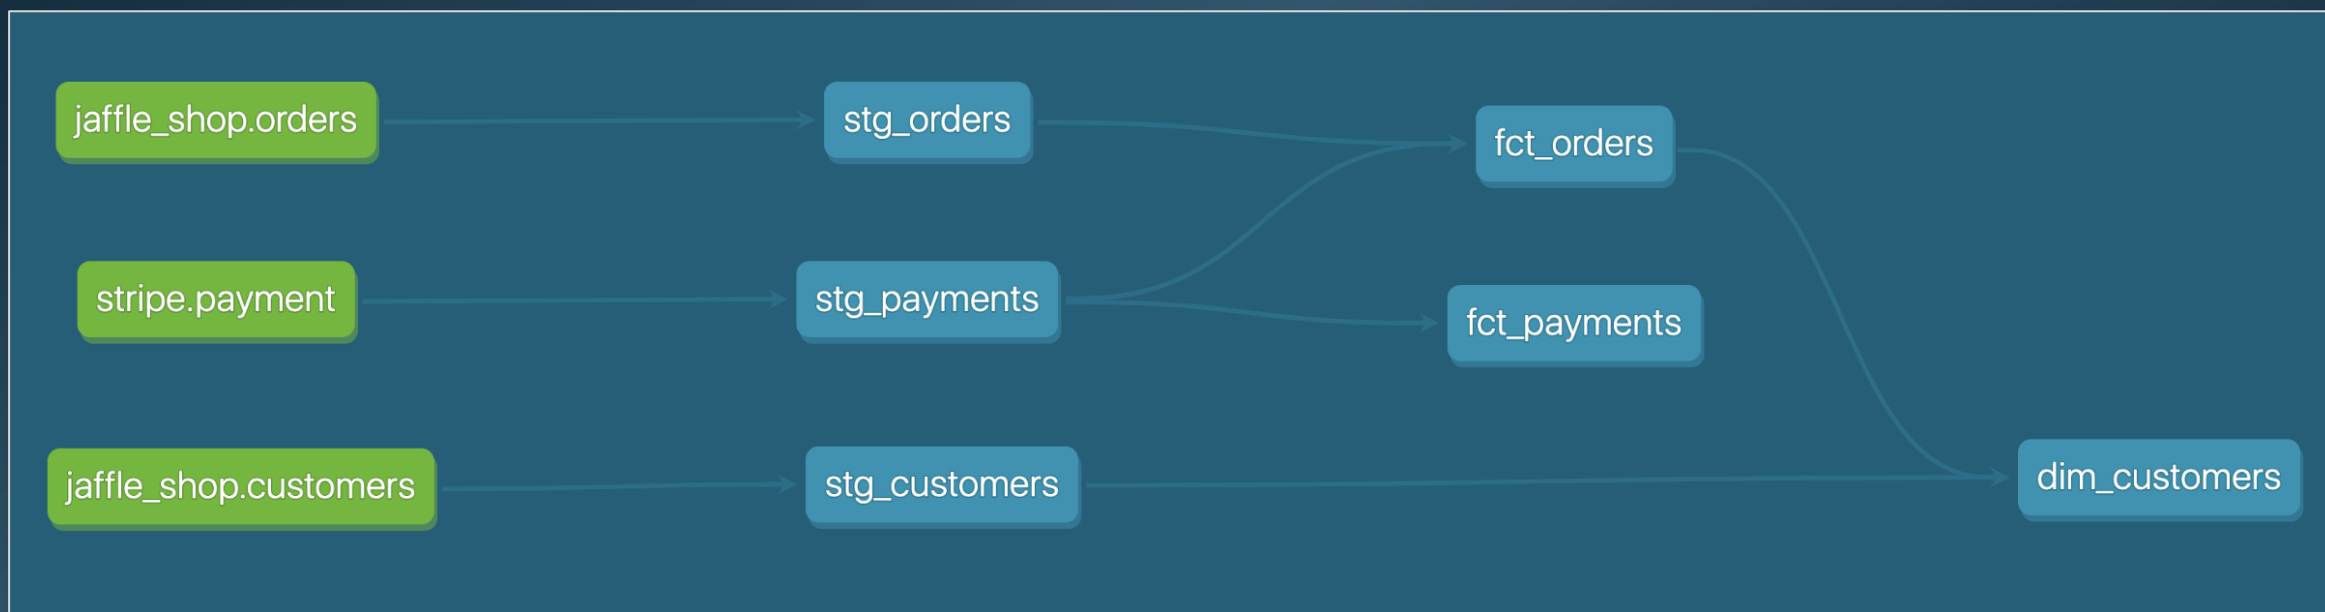


Data
Warehouse

dbt (including cloud) support in Warehouse

Declare your SQL loading processes

Create DAG for loading tables based on dependencies





Warehouse Utilization




Usage is tracked by Capacity Unit seconds (consumed by read and write activity against the Warehouse and reads against Lakehouse SQL Endpoint)

- Warehouse Query: Compute charged for Warehouse (includes user generated and system tasks)
- SQL Endpoint Query: Compute charged for Lakehouse SQL Endpoint (includes user generated and system tasks)
- One Lake compute: Compute charged for all reads and writes for data stored in One Lake

Operation name	CU (s)	Duration (s)	Users	Billing type
Warehouse Query	2,187.77	1,360.42	5	Billable
OneLake Compute	0.01	11,880.00	1	Billable
Total	2,187.78	13,240.42	6	

Operation name	CU (s)	Duration (s)	Users	Billing type
SQL Endpoint Query	4,086.87	1,503.81	6	Both
Total	4,086.87	1,503.81	6	



Bursting and Smoothing

- Regardless of SKU, Fabric *bursting* will automatically allocate resources as needed to execute at maximum performance
- As such, one query could consume all the quota of a single time window and much more!
- To avoid an overload, *smoothing* kicks in



Dynamics Management Views



Monitor Connection, Session, and Request Status in SQL Analytics Endpoint and Warehouses

`sys.dm_exec_connections`

- Provides comprehensive information about active connections to the Fabric Warehouse SQL Engine
- Includes details such as session ID, client address, client port, protocol, and authentication method.

`sys.dm_exec_sessions`

- shows information about all active user connections
- Offers insights into resource utilization by each session, such as the number of active requests

`sys.dm_exec_requests`

- Provides detailed information about currently executing or waiting queries
- Includes details such as session ID, query text, start time, status

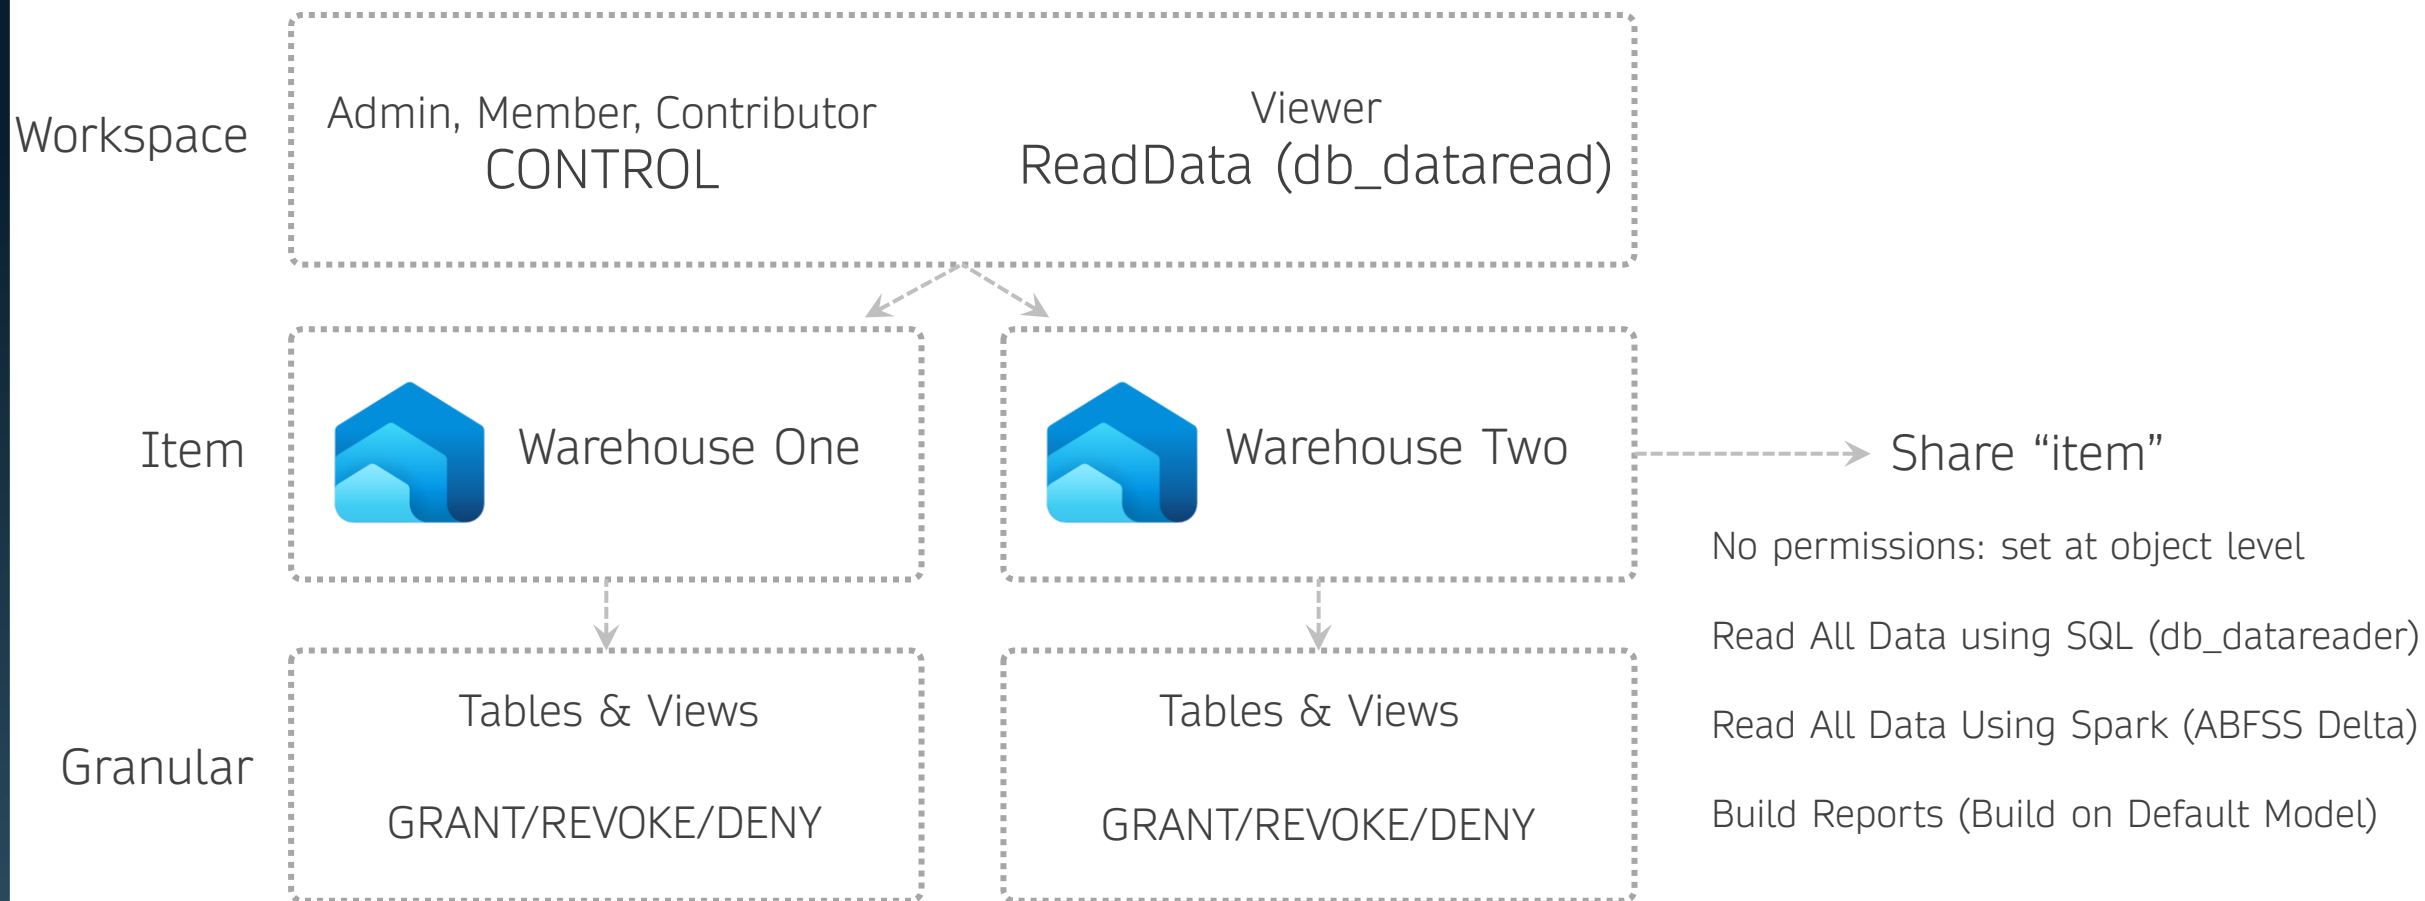
Member, Contributor, and Viewer can execute `sys.dm_exec_sessions` and `sys.dm_exec_requests`



Security



Security can be applied at Workspace, Warehouse, and Object level





Restore-In-Place

- Restore to a known “good” state
- Used when any corruption occurs
- Restore back after a failed deployment
- Restore back to a version for dev/test

Table Cloning (see next section)

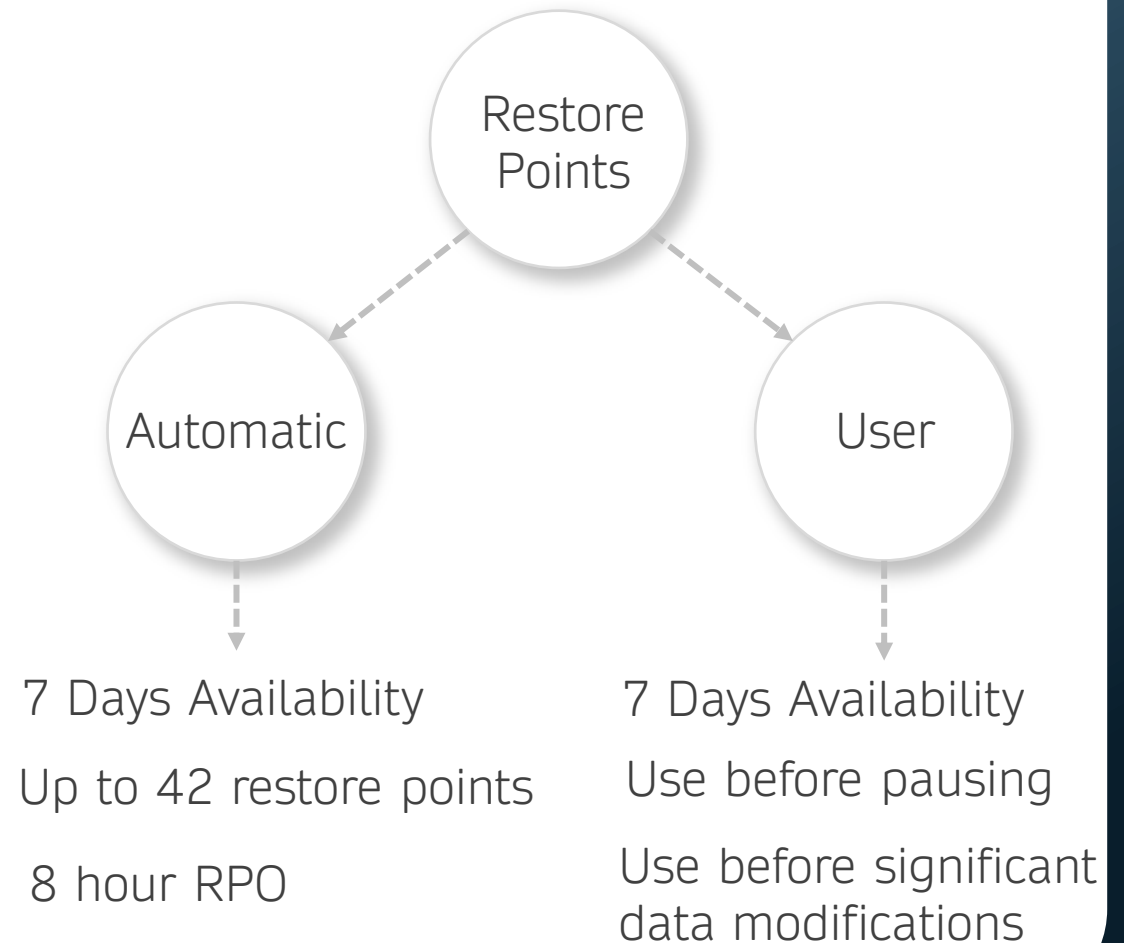




Table Cloning



Clone an existing table into new table

“Shallow” clone as only the metadata is cloned

Cloned table is separate from base table

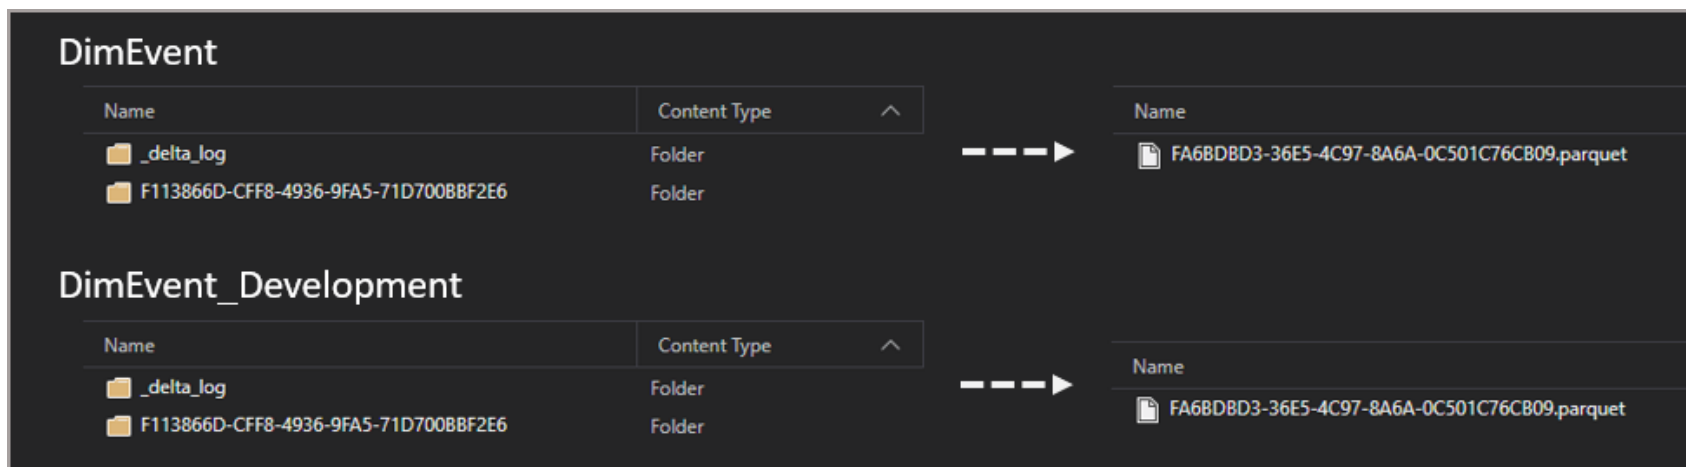
--create table clone

```
CREATE TABLE <schema>.<new_table_name>
```

```
AS CLONE OF <schema>.<existing_table_name>
```

Used for:

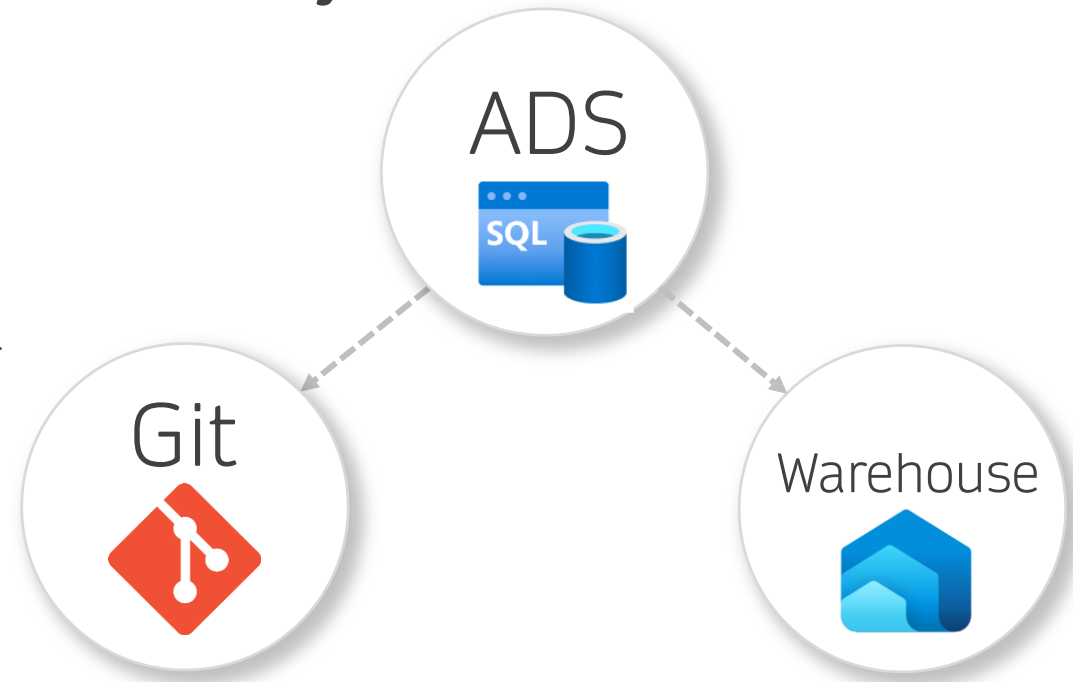
- Snapshotting at point-in-time
- Backups for recovery
- Testing data changes before applying to main table



“Treat your database like code” Grant Fritchey

Azure Data Studio Database Projects

- Current version 1.48.0 (DB: 1.4.2)
- Support for Warehouse and Lakehouse SQL Endpoint
- Uses dacpac for build and deploy
- Cannot ALTER!

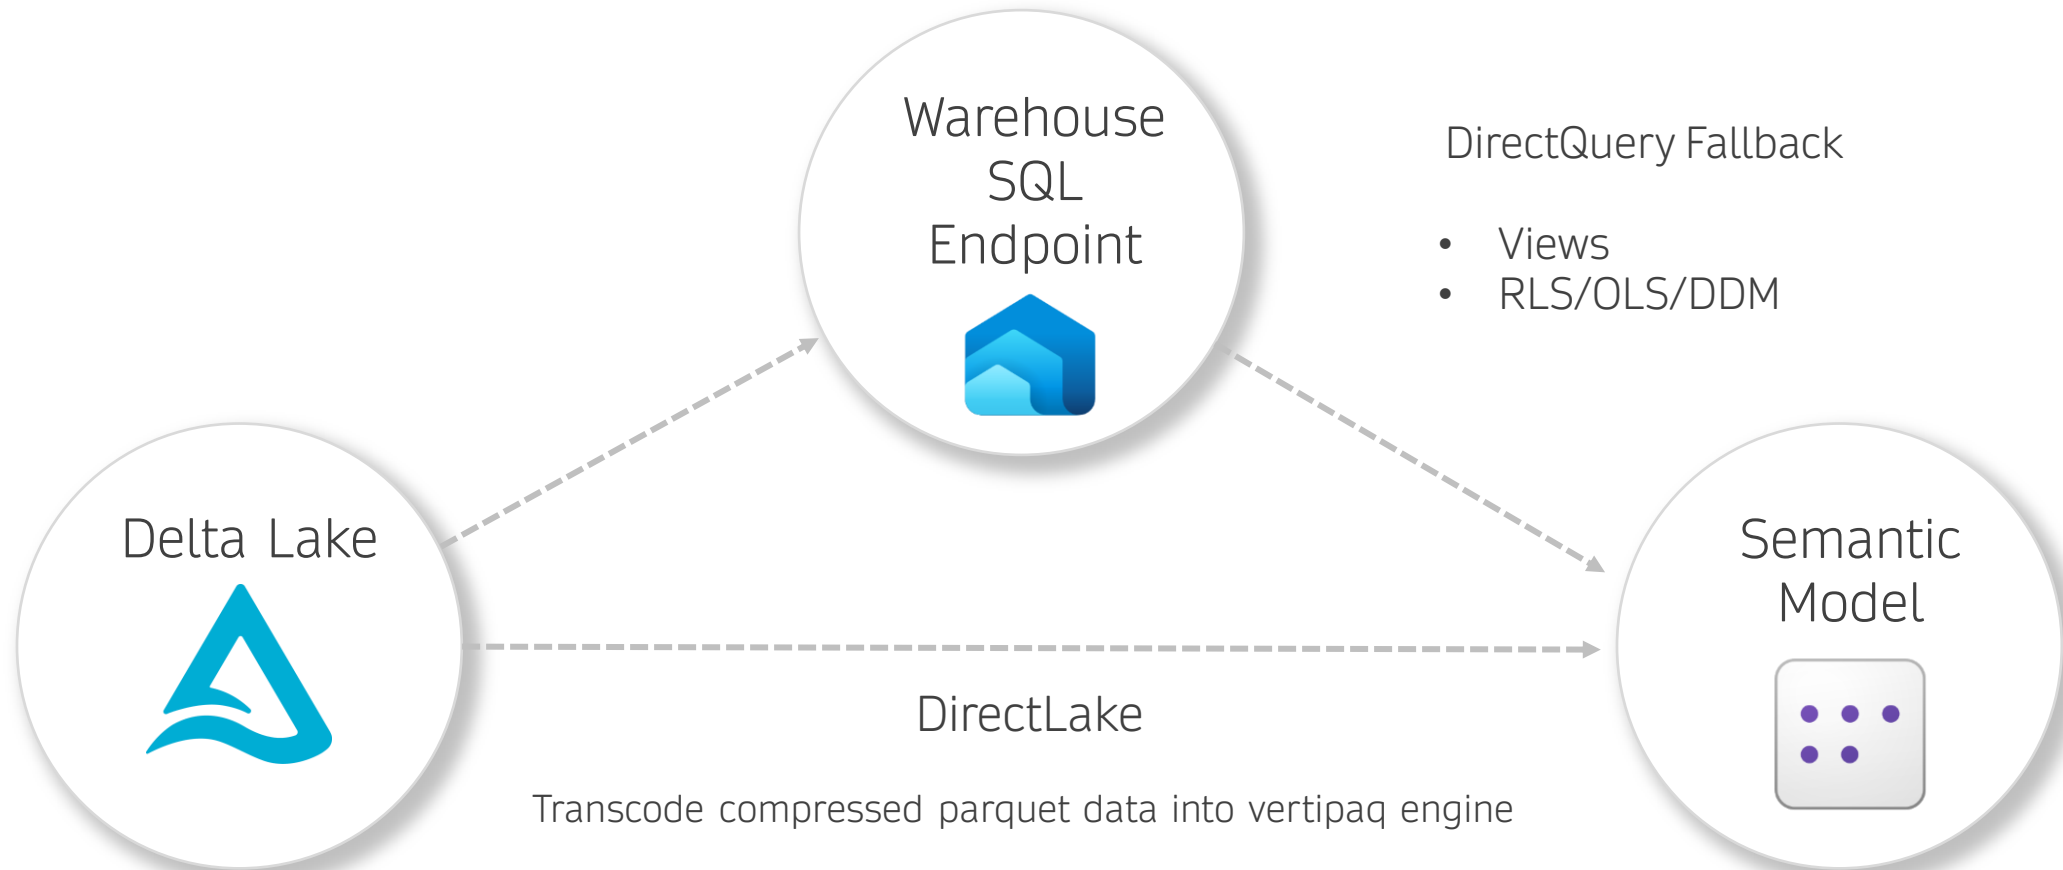


Connect ADS to Git to allow source control

Semantic Models – Direct Lake

Direct Lake is the new connectivity method from Semantic Models to Warehouse & Lakehouse

No requirement to Import the data into the Semantic model



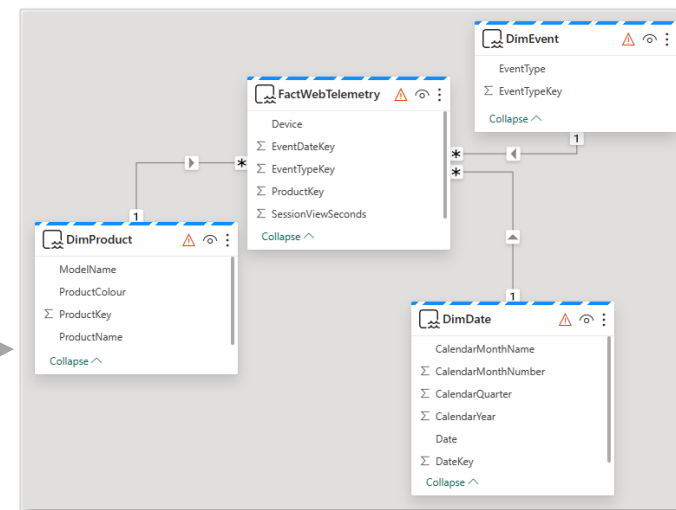


Constraints created in Warehouse appear in Default Semantic Model

Constraints created in Default Semantic Model appear in Warehouse

	ABC foreign_table	ABC primary_table	ABC fk_column	ABC pk_column ↑	ABC fk_constraint
1	dbo.FactWebTelemetry	dbo.DimDate	EventDateKey	DateKey	FK_FactWebTelemetry_EventDateKey
2	dbo.FactWebTelemetry	dbo.DimEvent	EventTypeKey	EventTypeKey	FK_FactWebTelemetry_EventTypeKey
3	dbo.FactWebTelemetry	dbo.DimProduct	ProductKey	ProductKey	FK_FactWebTelemetry_ProductKey

Default Semantic Model has a bi-directional sync with it's underlying Warehouse



Custom Semantic Models are unaffected by Warehouse constraints

A black and white portrait of a man with a beard and mustache, wearing a dark shirt and a light-colored vest, looking towards the camera.

Thank you

Connect with me at:

 <https://linkedin.com/in/brianbonk>
 <https://brianbonk.dk>
 <https://github.com/brianbonk>

Connect

