




DP-203T00: Ingest and load data into the data warehouse



Agenda



Lesson 01: Use data loading best practices in Azure Synapse Analytics

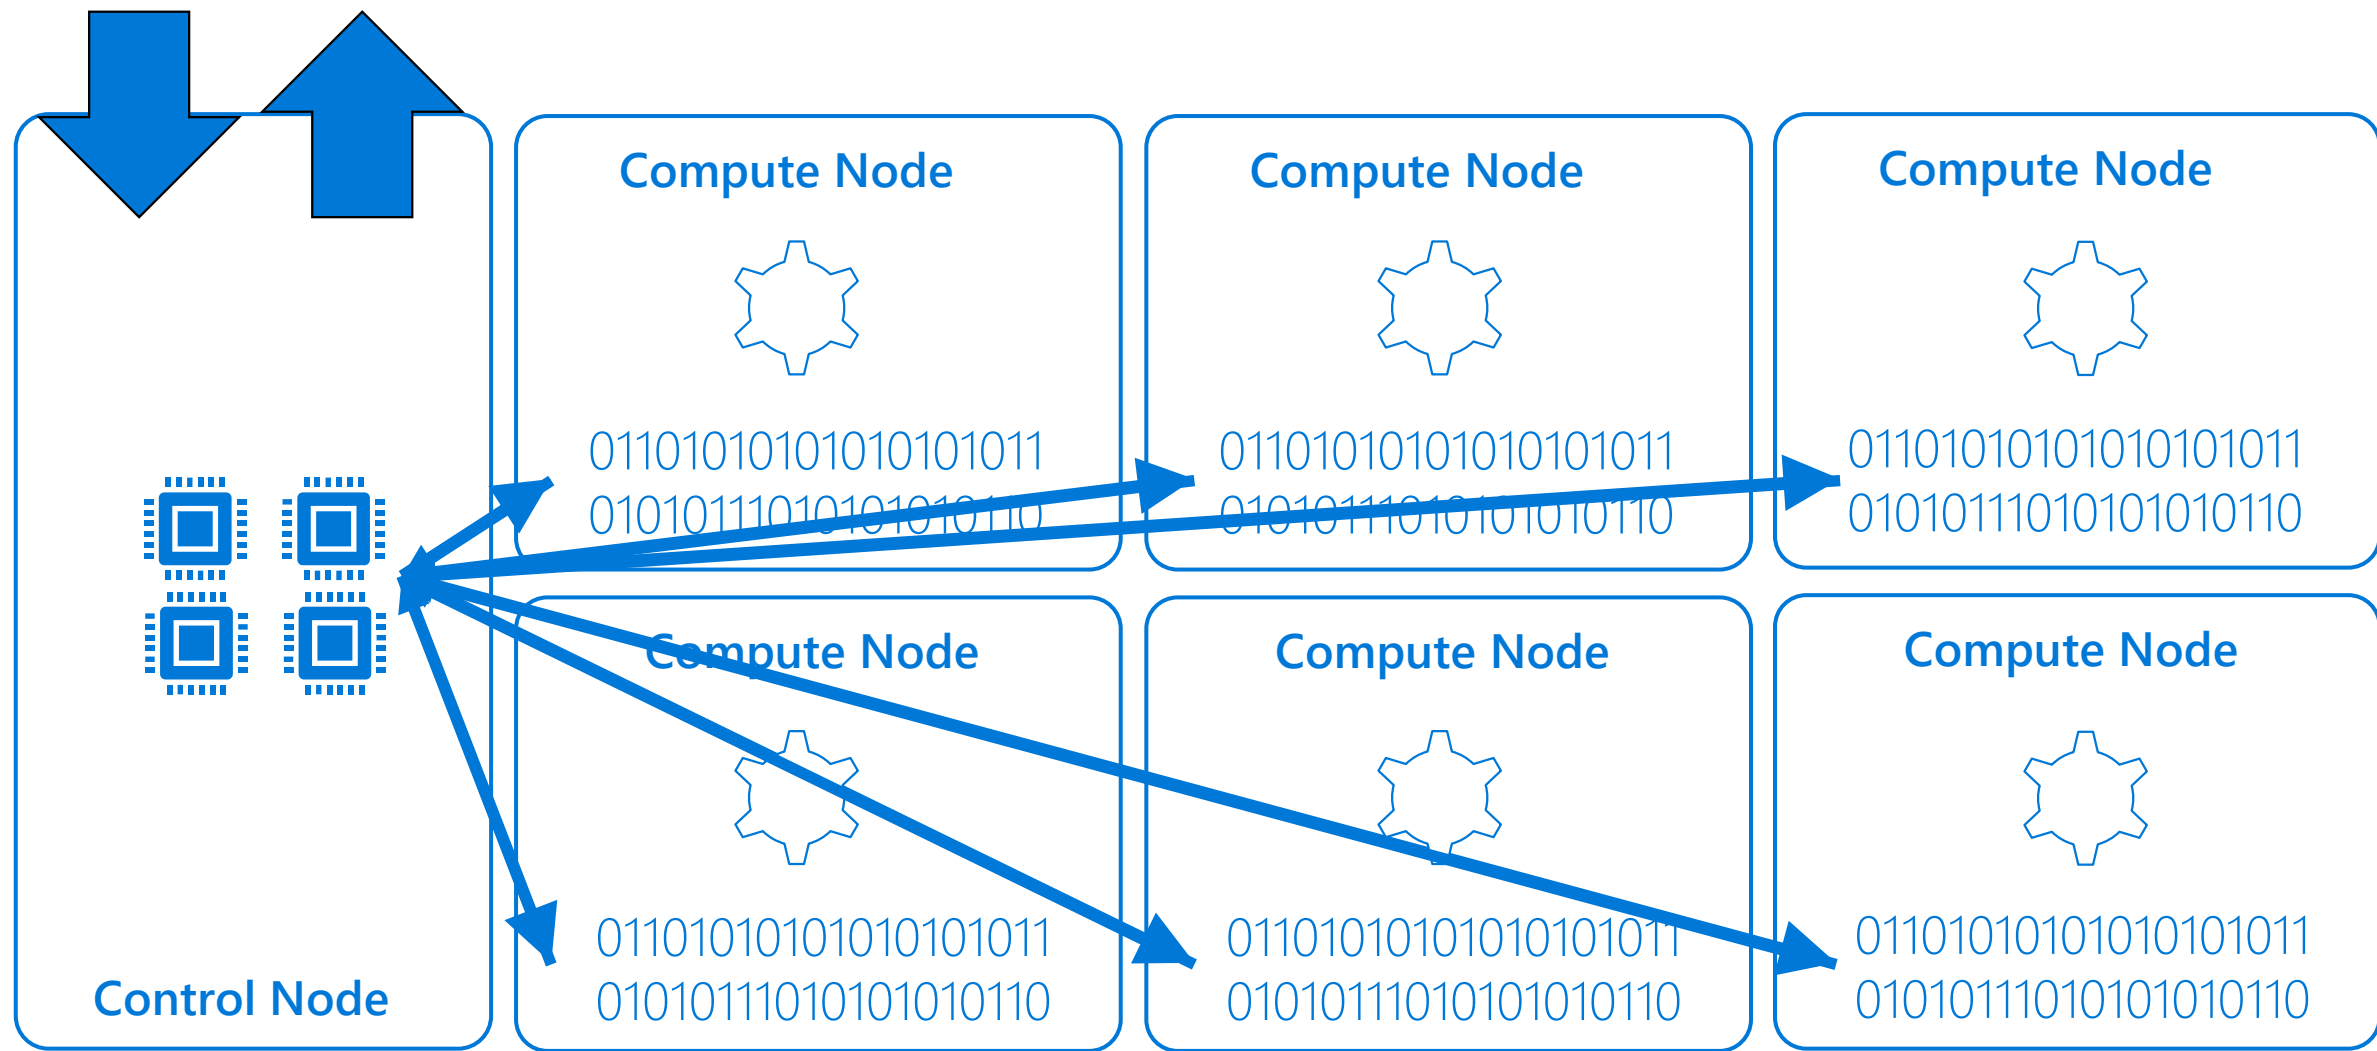


Lesson 02: Petabyte-scale ingestion with Azure Data Factory

Lesson 01: Use data loading best practices in Azure Synapse Analytics



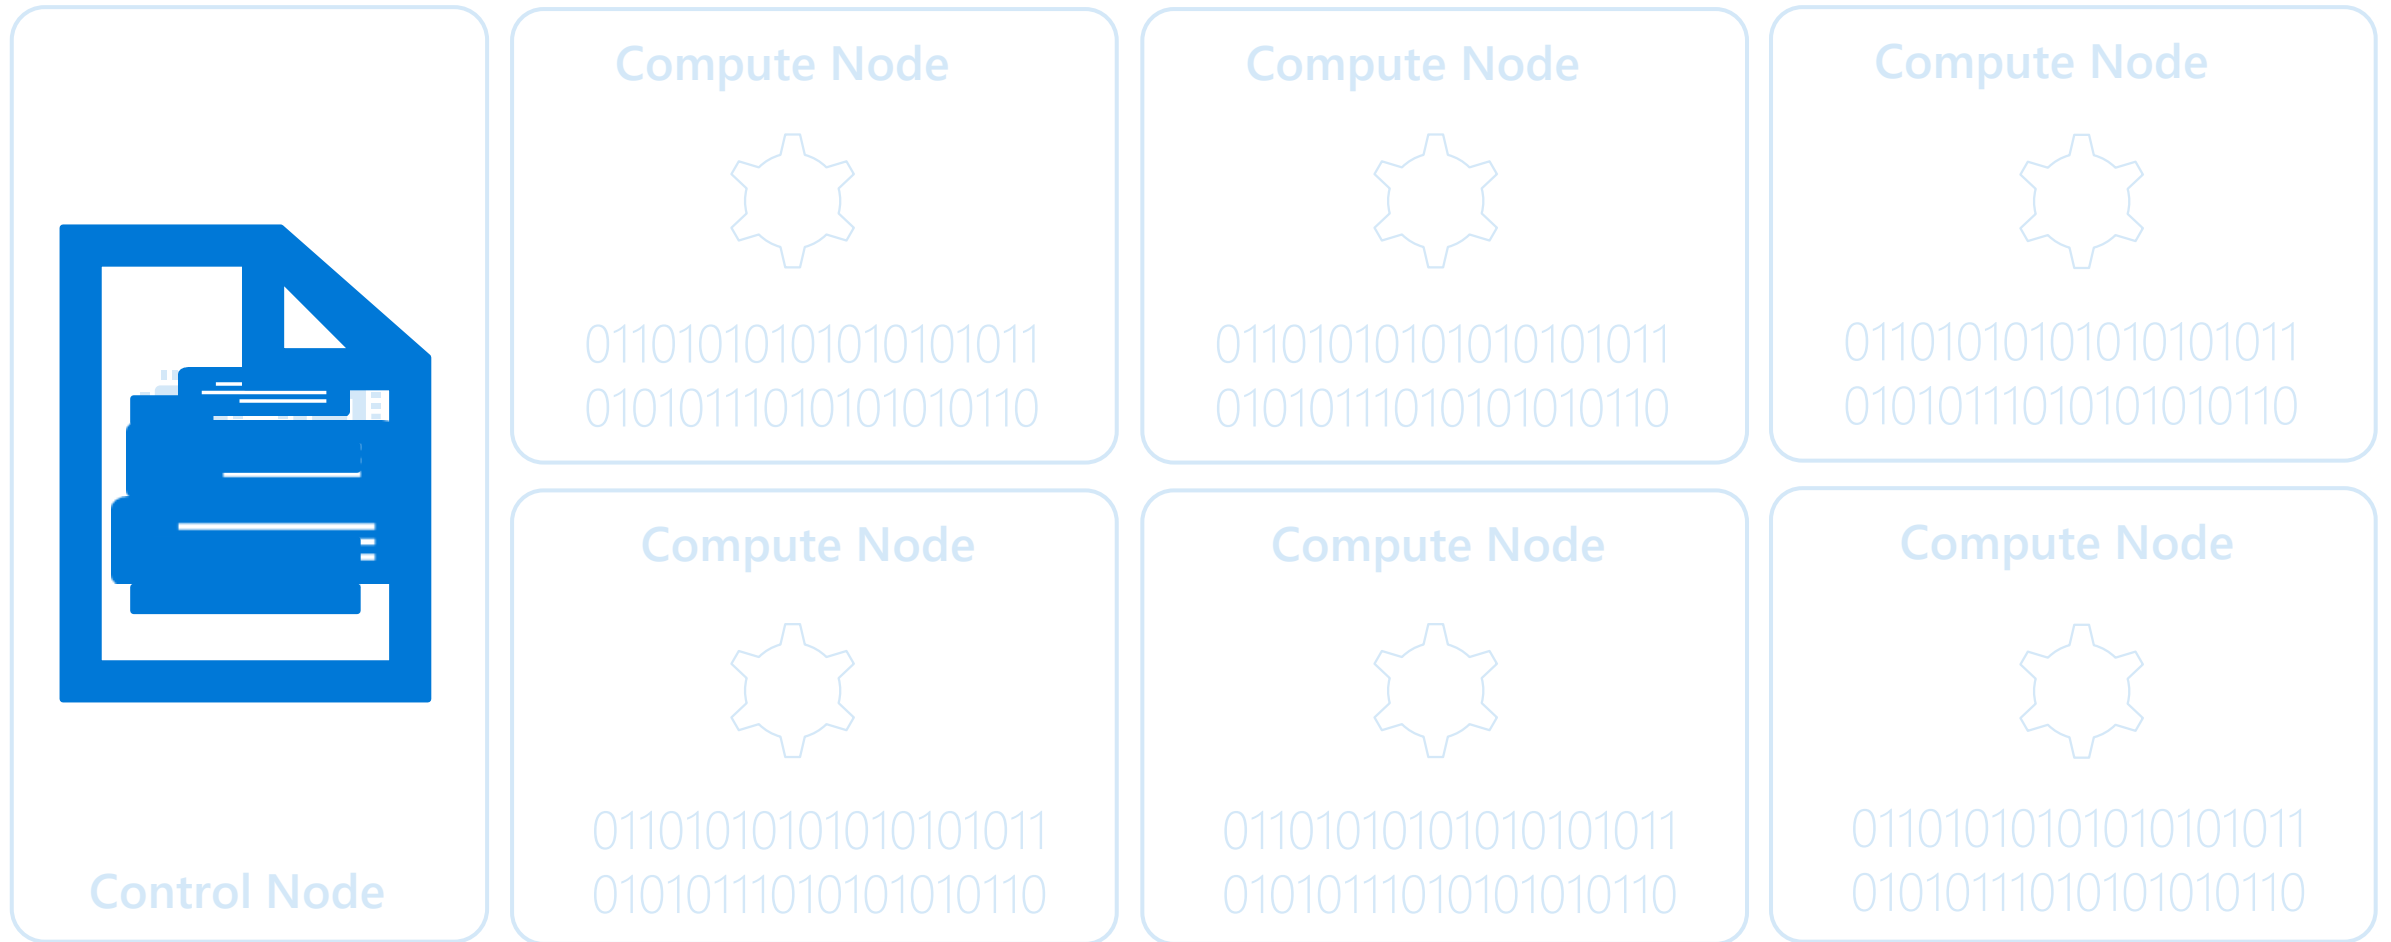
Dedicated SQL Pool architecture revision



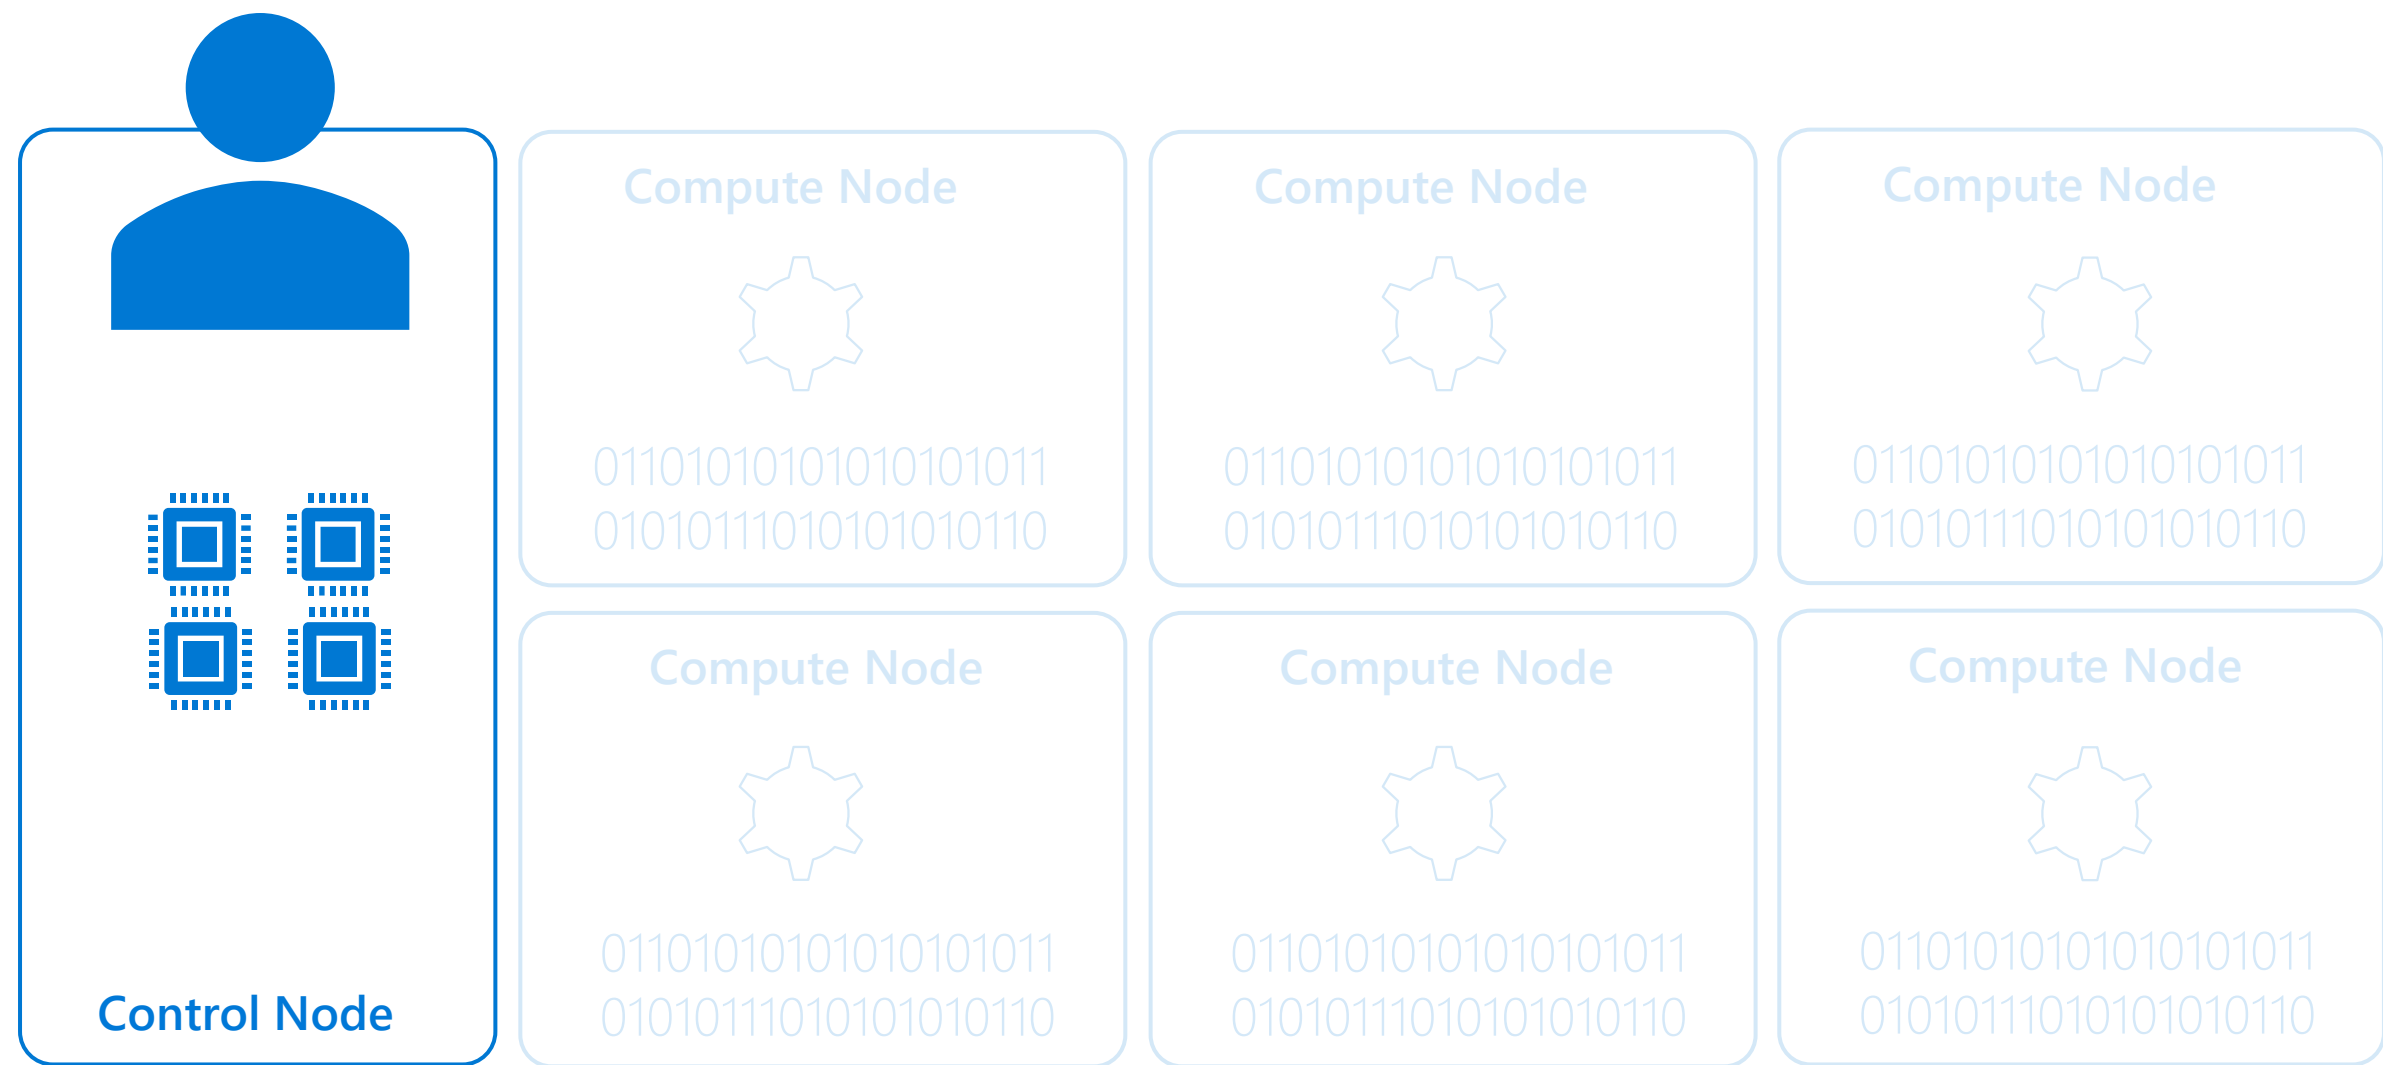
Understand data load design goals

- Where is my data coming from?
- Is the data nett new? or do you receive changes from existing datasets?
- How often is the data being refreshed, added to or replaced?
- What formats are the data coming in?
- Is the data ingestible as-is? or are transformations and cleansing tasks required?
- Which takes priority, loading or querying/analysis?

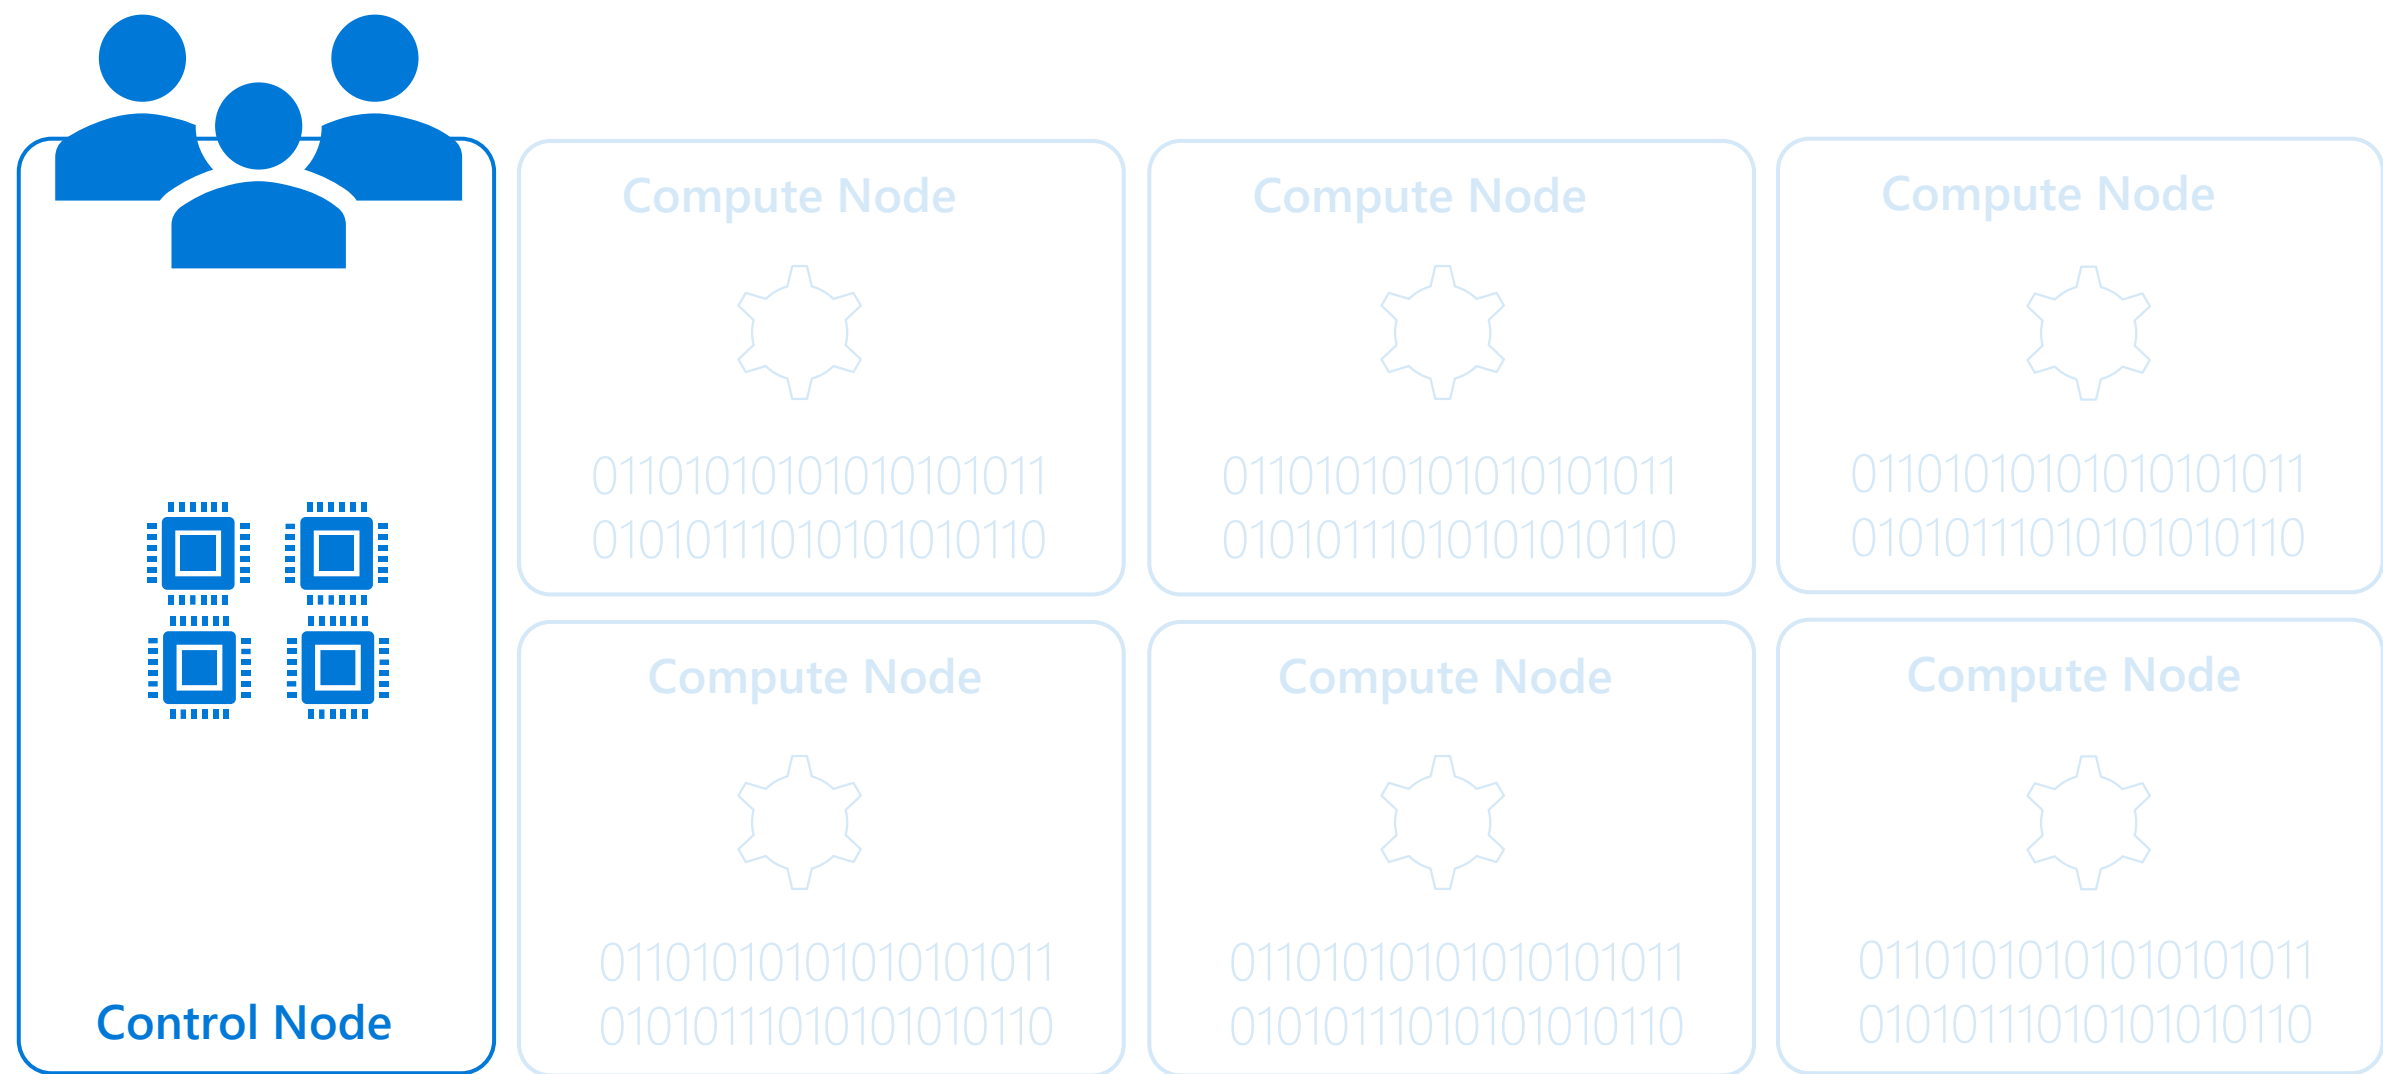
Manage singleton updates



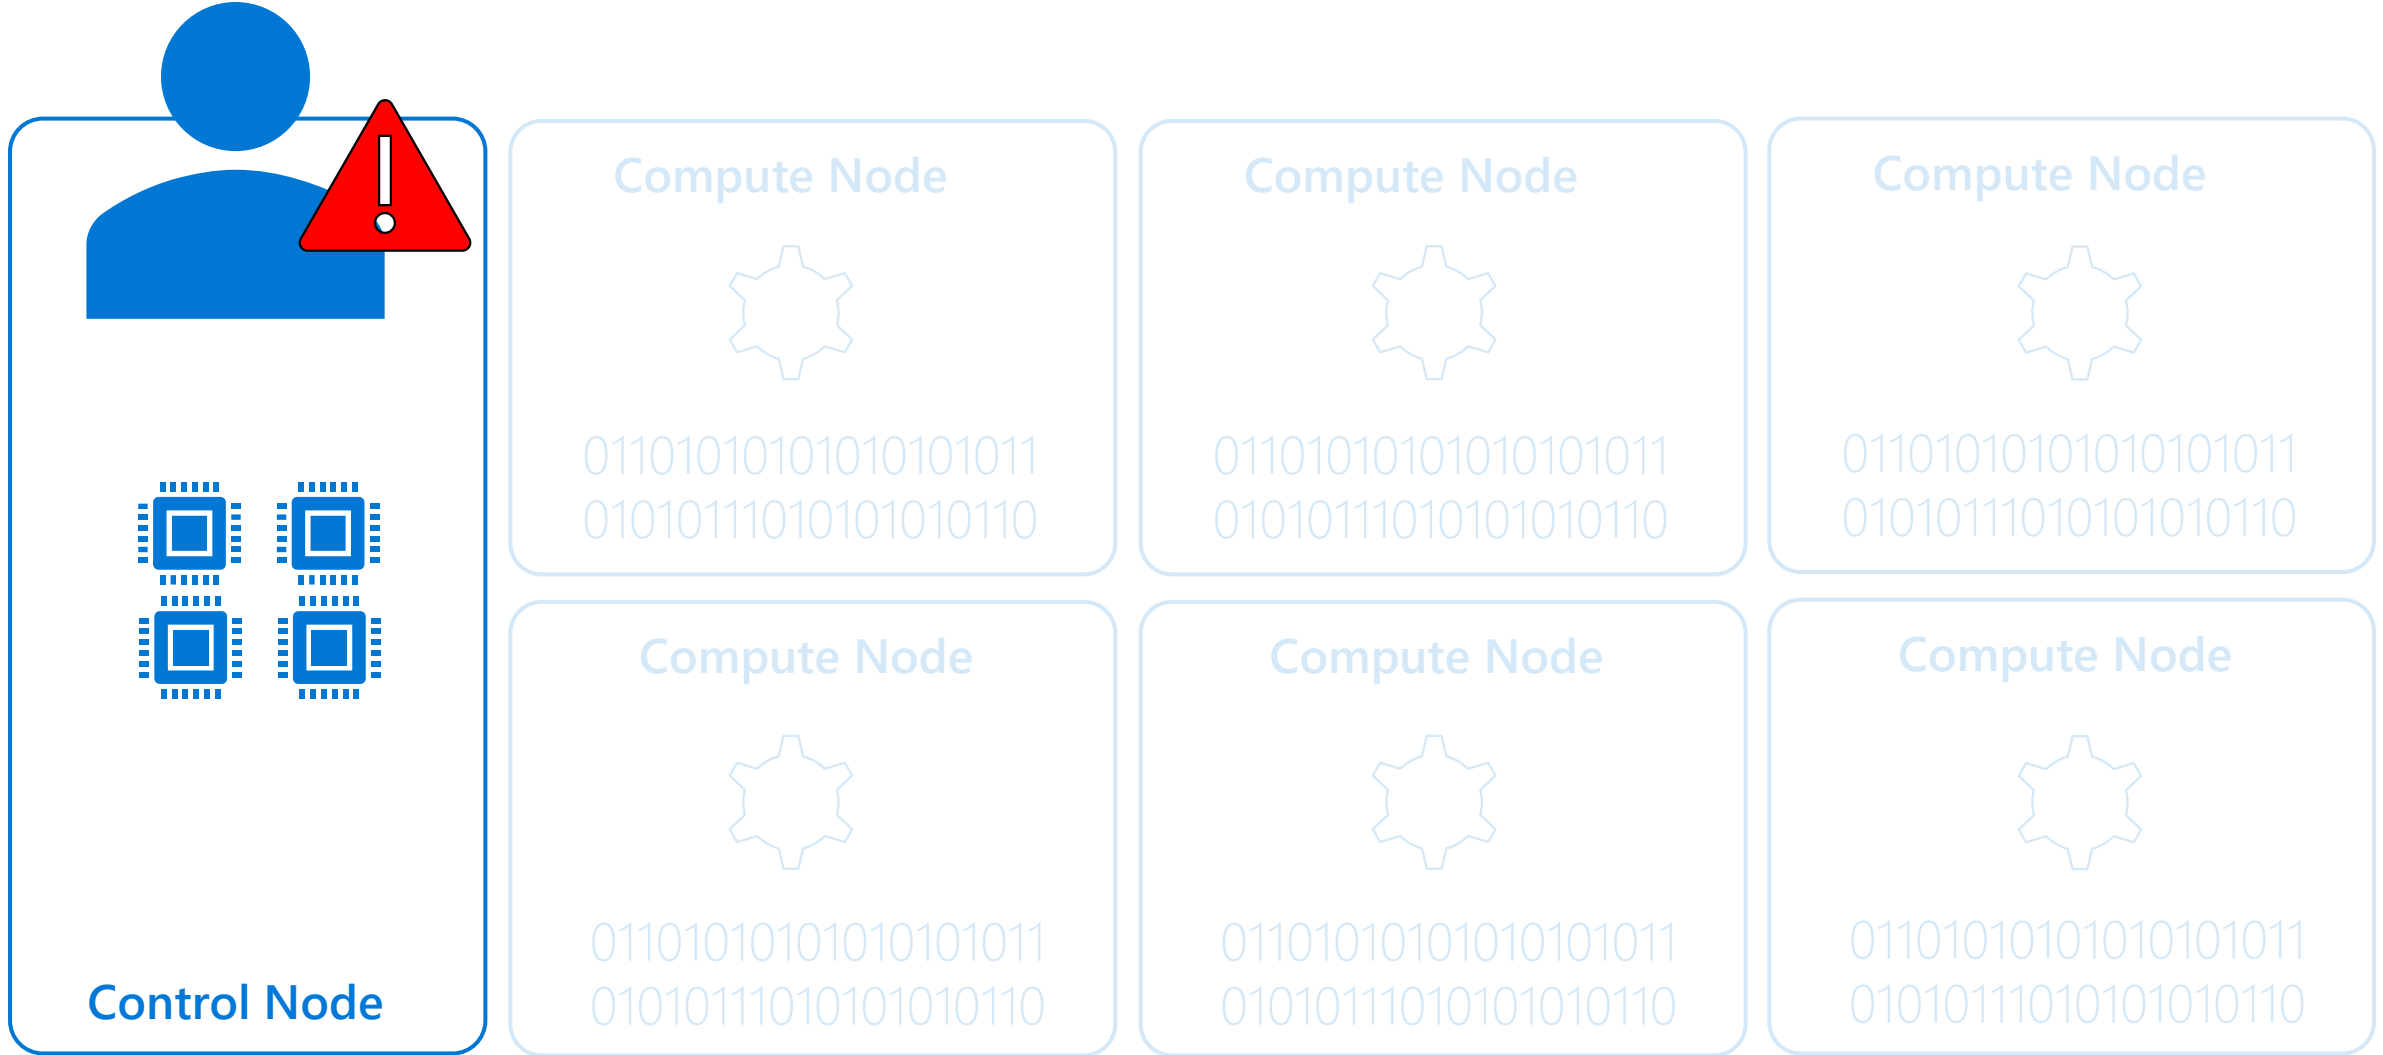
Set-up dedicated data loading accounts



Manage concurrent access to Azure Synapse Analytics




The diagram illustrates a 1U server rack configuration. On the left, a large blue box represents the **Control Node**, which contains a blue server icon with a red warning triangle and four blue chip icons. To the right of the Control Node are six smaller, light blue boxes, each representing a **Compute Node**. Each Compute Node contains a gear icon and two lines of binary code: 01101010101010101011 and 01010111010101010110.

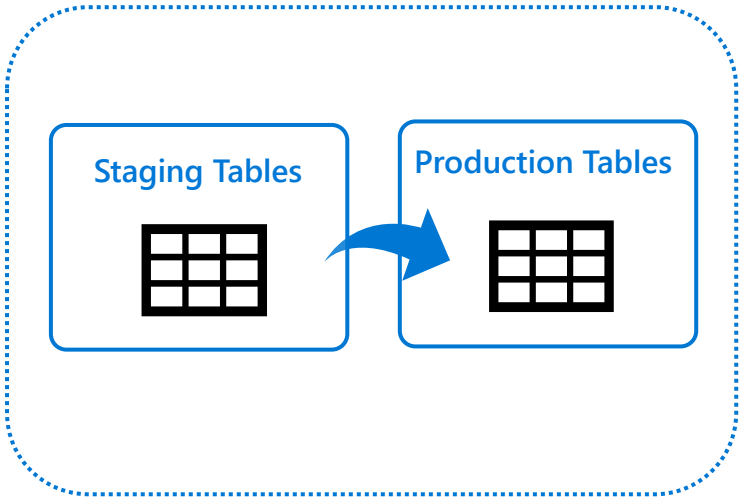


Use PolyBase, the Copy command or the Copy Activity


PolyBase



Copy command



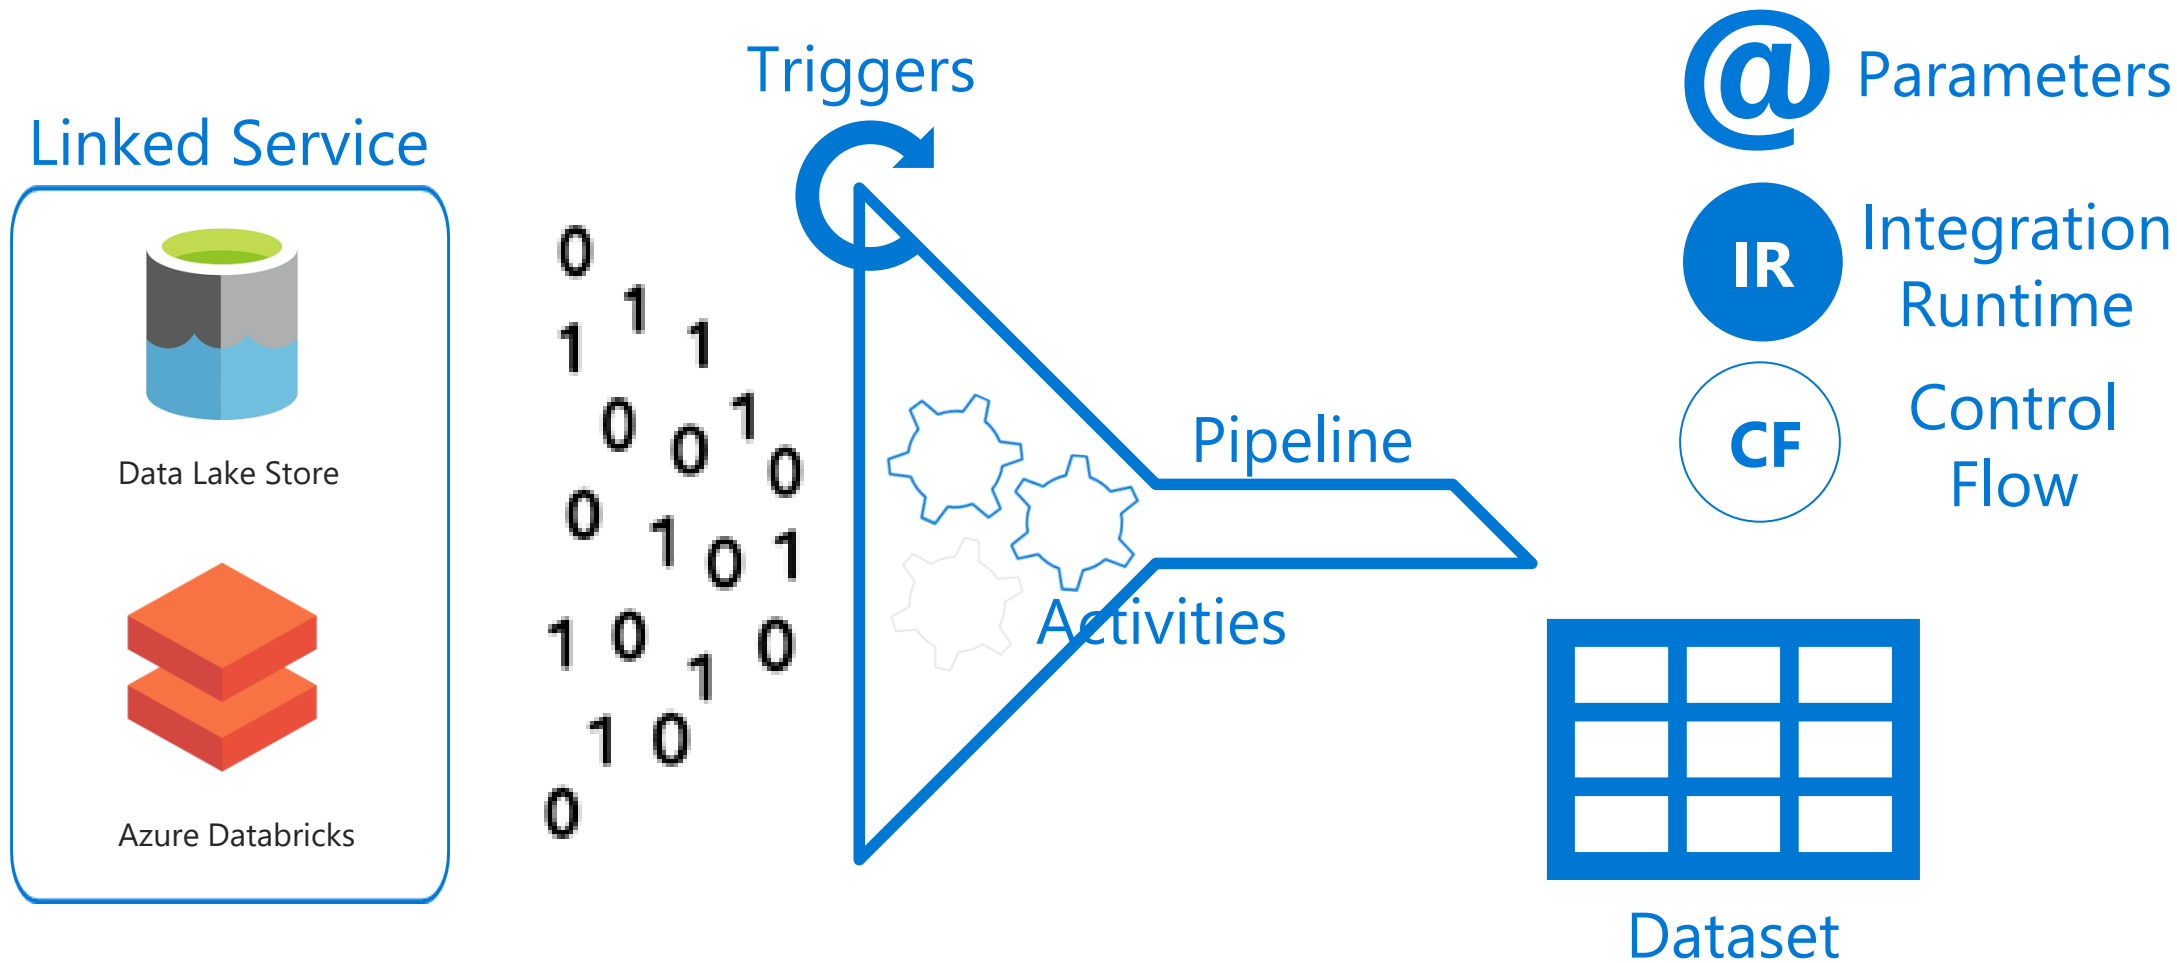
Copy data activity




Lesson 02: Petabyte-scale ingestion with Azure Data Factory





Azure Data Factory/Synapse pipeline revision




Petabyte-scale ingestion with Azure Data Factory



Understanding integration




Review questions




Q01 – Which data loading feature limits the number of resources a group of requests can consume in Azure Synapse Analytics?

A01 – [Workload management](#)



Q02 – Why should you split up one large files into smaller files when loading data into a dedicated SQL pool in Azure Synapse Analytics?


A02 – [To take advantage of the Massively Parallel Processing \(MPP\) architecture](#)



Q03 – In which section of the Data Factory designer canvass would you find the Copy Activity?

A03 – [Move and Transform](#)

Lab: Ingest and load data into the data warehouse



Lab overview

This lab teaches students how to ingest data into the data warehouse through T-SQL scripts and Synapse Analytics integration pipelines. The student will learn how to load data into Synapse dedicated SQL pools with PolyBase and COPY using T-SQL. The student will also learn how to use workload management along with a Copy activity in a Azure Synapse pipeline for petabyte-scale data ingestion.


Lab objectives

After completing this lab, you will be able to:


Use data loading best practices in Azure Synapse Analytics

Petabyte-scale ingestion with Azure Data Factory


Lab review




Question 1 – What is the appropriate distribution to use in a staging table?



Question 2 – When creating an External Data Source, what should the Type property be set to access data in a data lake?



Question 3 – When setting file format properties for a data source in the Copy Activity, you can import the schema from a connection/store and what other option?



Question 4 – You want to truncate a staging table before loading it with data in a Copy activity. Which sink property would you set to do this?

Module summary

In this module, you have learned about:

Use data loading best practices in Azure Synapse Analytics

Petabyte-scale ingestion with Azure Data Factory

Next steps

After the course, consider reading [[Best practices for loading data using dedicated SQL pools in Azure Synapse Analytics](#)] for more guidance on data ingestion and loading

