

**Fellowwind**

# Real-Time Analytics at scale in Fabric

A unified analytics solution for the era of AI



redgate

TIMEXTENDER






# Brian Bønk Rueløkke

Principal & Enterprise architect, Data & AI

*Fellowmind*

 <https://linkedin.com/in/brianbonk>  
 <https://brianbonk.dk>  
 <https://github.com/brianbonk>





# Jaques Cousteau

## 1910-1997



# Jaques Cousteau

## 1910-1997



# The history of Kusto



Azure Sentinel



Log Analytics



Real-Time Analytics



Azure resource  
graph




Microsoft 365  
Defender

CMPIvot

CMPIvot





The world is exploding with  
high-granularity data...





It all starts with data



Telemetry – a key data for digital transformation



Telemetry – a key data for digital transformation





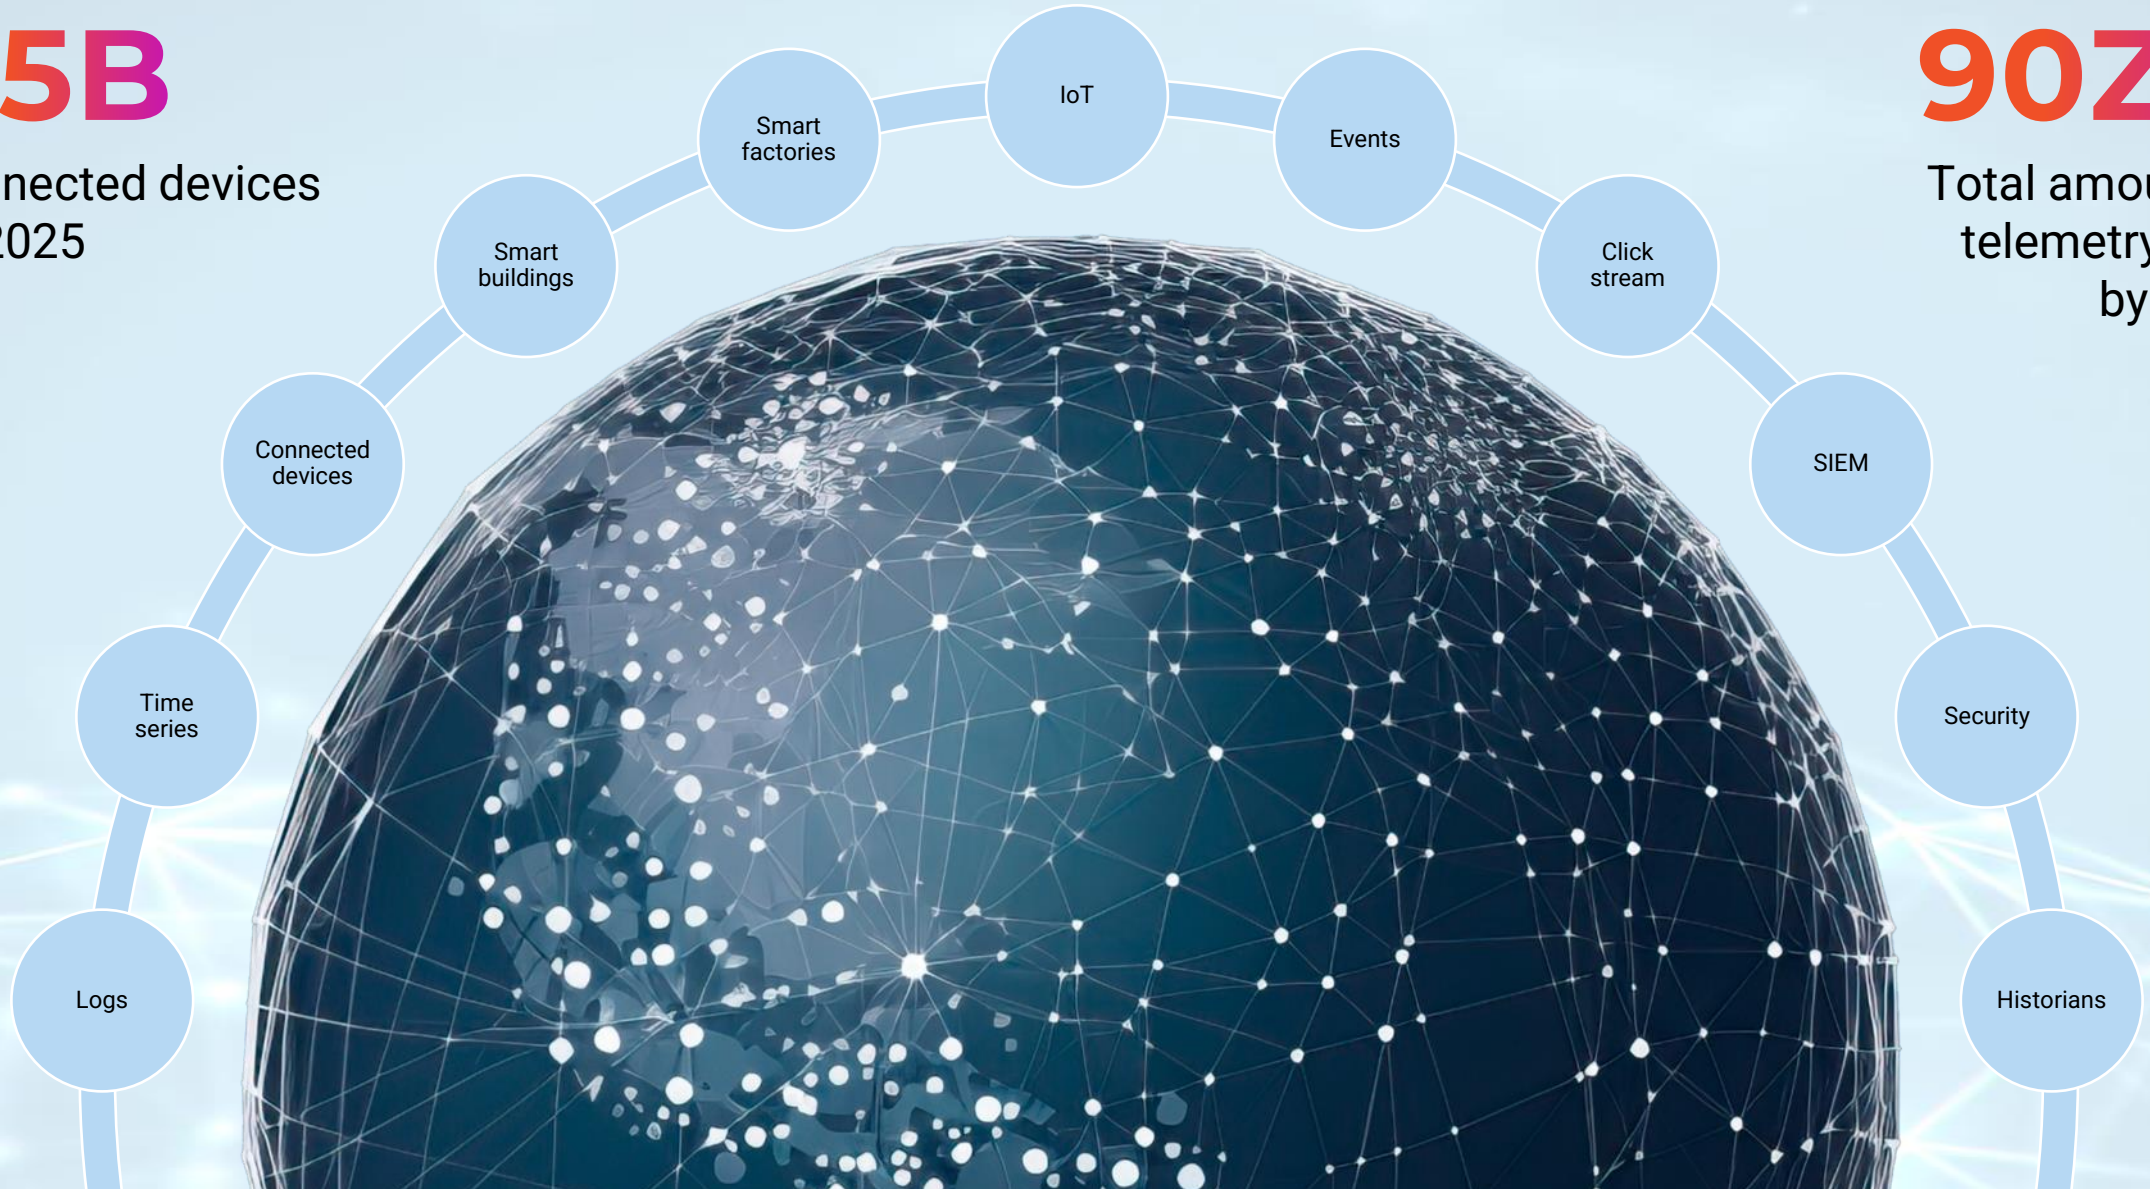
# Telemetry – a key data for digital transformation

**75B**

Connected devices  
by 2025

**90ZB**

Total amount of  
telemetry data  
by 2025





# Digital transformation

Cybersecurity  
Asset tracking and management  
Predictive maintenance  
Supply chain optimization  
Customer experience  
Energy management  
Inventory management  
Quality control  
Environmental monitoring  
Fleet management  
Health and safety





# Microsoft Fabric

## Get data



Data Factory

## Prepare data



Synapse Data Engineering



Synapse Data Warehouse



Synapse Data Science



Synapse Real-Time Analytics

## Use data



Power BI



Data Activator

## Store data



OneLake



# Microsoft Fabric

## Get data



Data Factory

## Prepare data



Synapse Data Engineering



Synapse Data Warehouse



Synapse Data Science



Synapse Real-Time Analytics

## Use data



Power BI



Data Activator

## Store data



OneLake



**Fabric Real-time Analytics** solution enables organizations to consume **vast amount of data**, focus and **scale up** their Analytics solution with **data in motion**, **empower** their business analysts, and **democratize their data** for citizen data scientists and Data Engineers




Synapse Real-Time  
Analytics

# ⚡ Streaming data with ease



## EVENTSTREAM

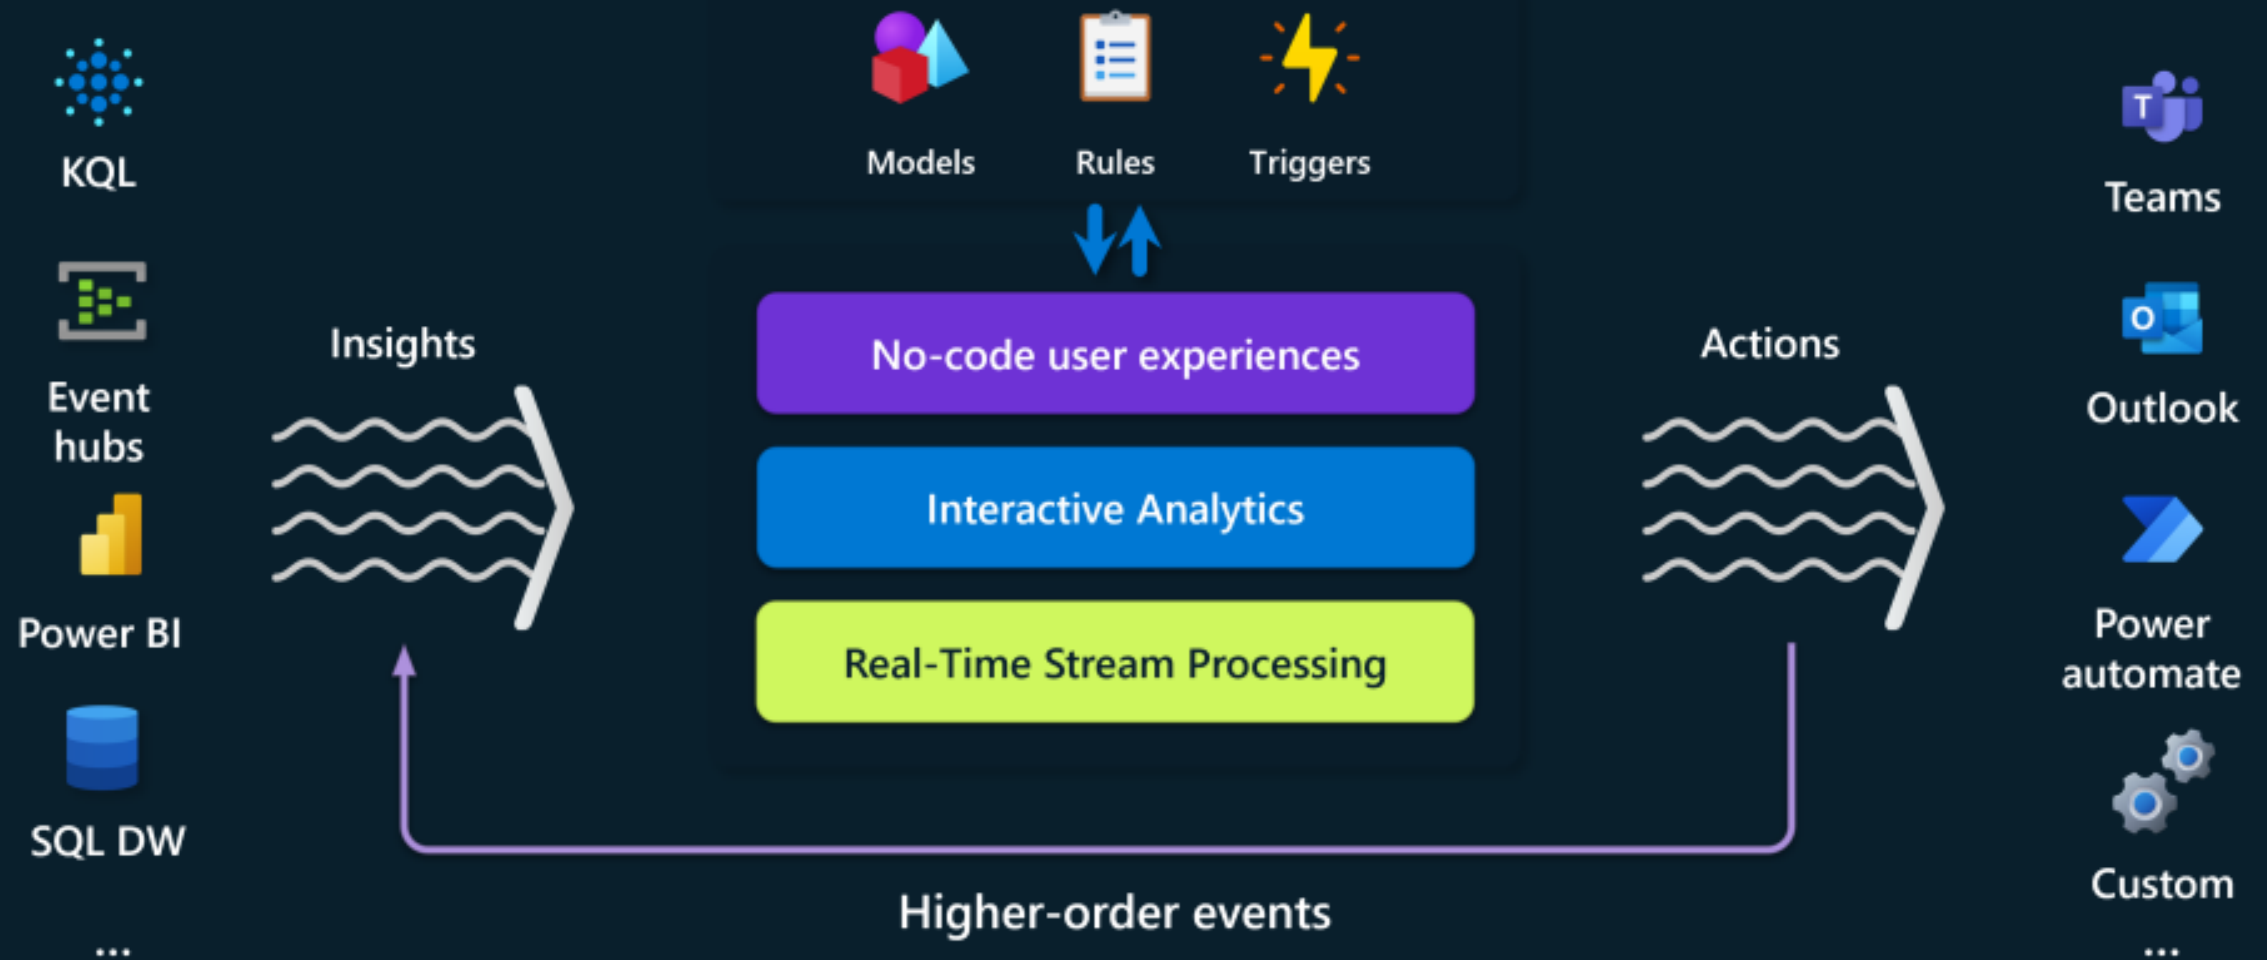
The brand-new event stream service, leverages the ability to get data from several sources of streaming data and save it to a wide variety of destinations, including OneLake, KQL databases and Azure services.



The service computes the data once and can pipe it out to several destinations at once. All configured and maintained from within the Microsoft Fabric portal and “coded” with your mouse.

Imagine scenarios of IoT devices loading data to both the data warehouse and other 3-rd party destinations – this can now be done using the low-code approach from Event Stream.

# Data Activator





# KQL database

## Key capabilities

Unlimited Scale  
(query, ingestion  
and storage)

Any data source

Any data format

Structured  
Semi-structured  
Free-text

Real-time  
transformation of  
complicated data  
structures

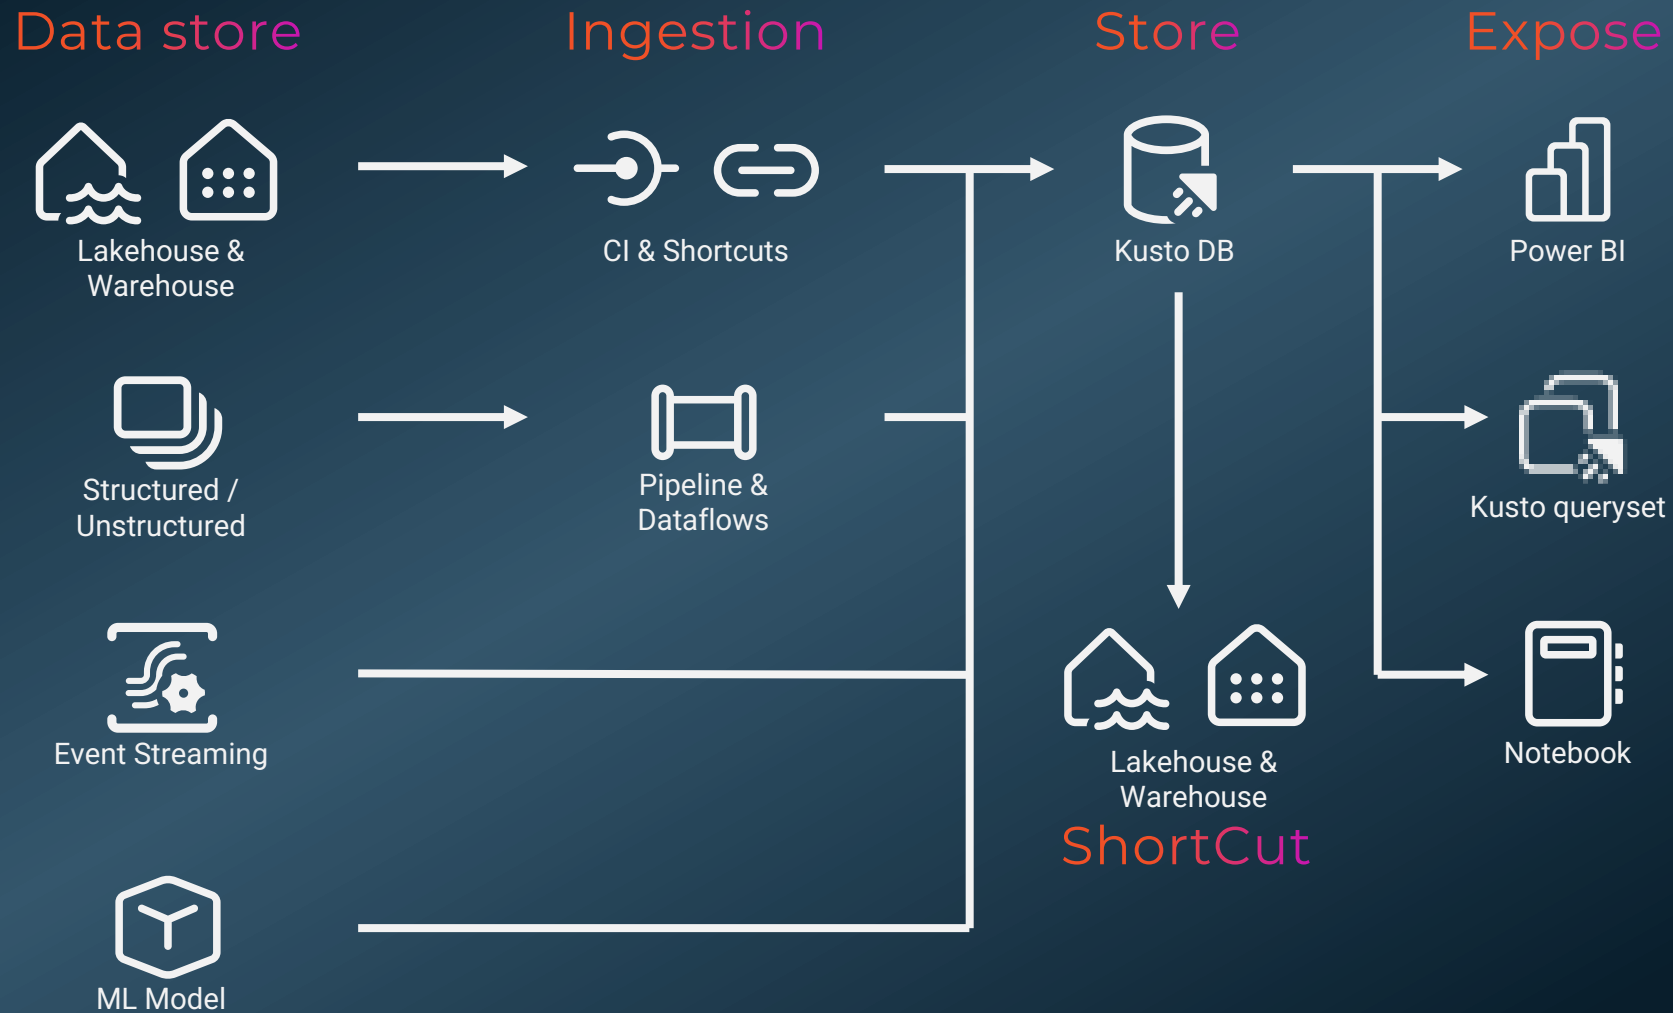
Streaming analytics in  
Near-Real-Time

High performance  
Low latency  
High freshness

Timeseries database

Everything is indexed  
and partitioned


# Real-Time Analytics



Get started for free


<https://dataexplorer.azure.com/freecluster>

<https://detective.kusto.io>



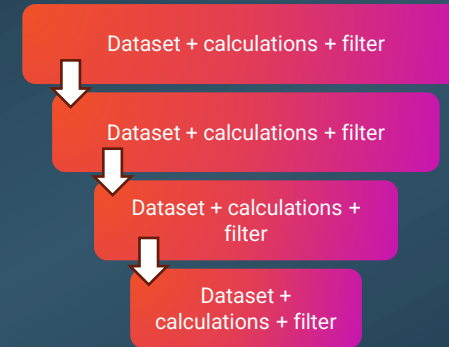


# The language and structure



# The language and structure

KQL



NYCTaxi

```
| where passenger_count > 1  
| project passenger_count, total_amount, VendorID, fare_amount  
| extend AmtPsngr = total_amount / passenger_count  
| where AmtPsngr > 10  
| summarize TotalAmount = sum(total_amount), AvgAmtPsngr = avg(AmtPsngr) by VendorID  
| where VendorID <> 1
```

Kusto in Power BI

Forget everything you know about  
query performance vs data types  
&  
data modelling best practices



# Data modelling Kusto in Power BI

- Single table reporting can be a good option, if you can include all columns from dimensions to the table
- M:M relations are hard to avoid, but not a big deal → all queries will be translated to KQL
- All dimensions must be tagged with “IsDimension=true”
- Dimensions can be imported if they are <1 mio rows.
- INTEGER and DECIMAL er slow joins compared to STRING



# Harness the Power (BI) of Kusto

Let Power BI build the KQL

- In Power Query
- Using DAX

Or build a Kusto  
function





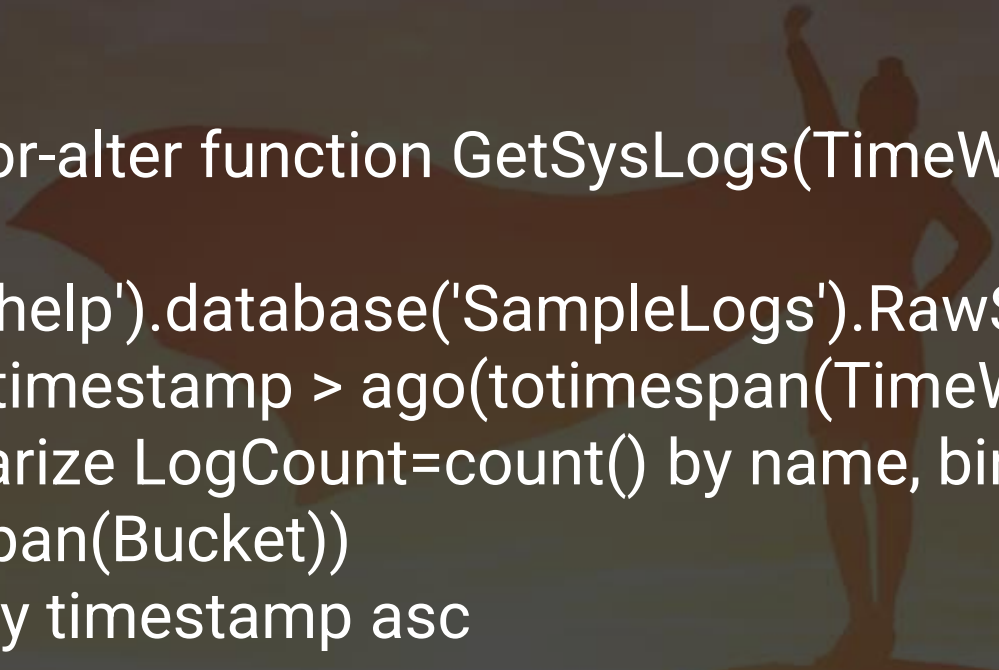


# DEMO

Live coding  
(hopefully no demo-ghost 🐻)



# Harness the Power (BI) of Kusto



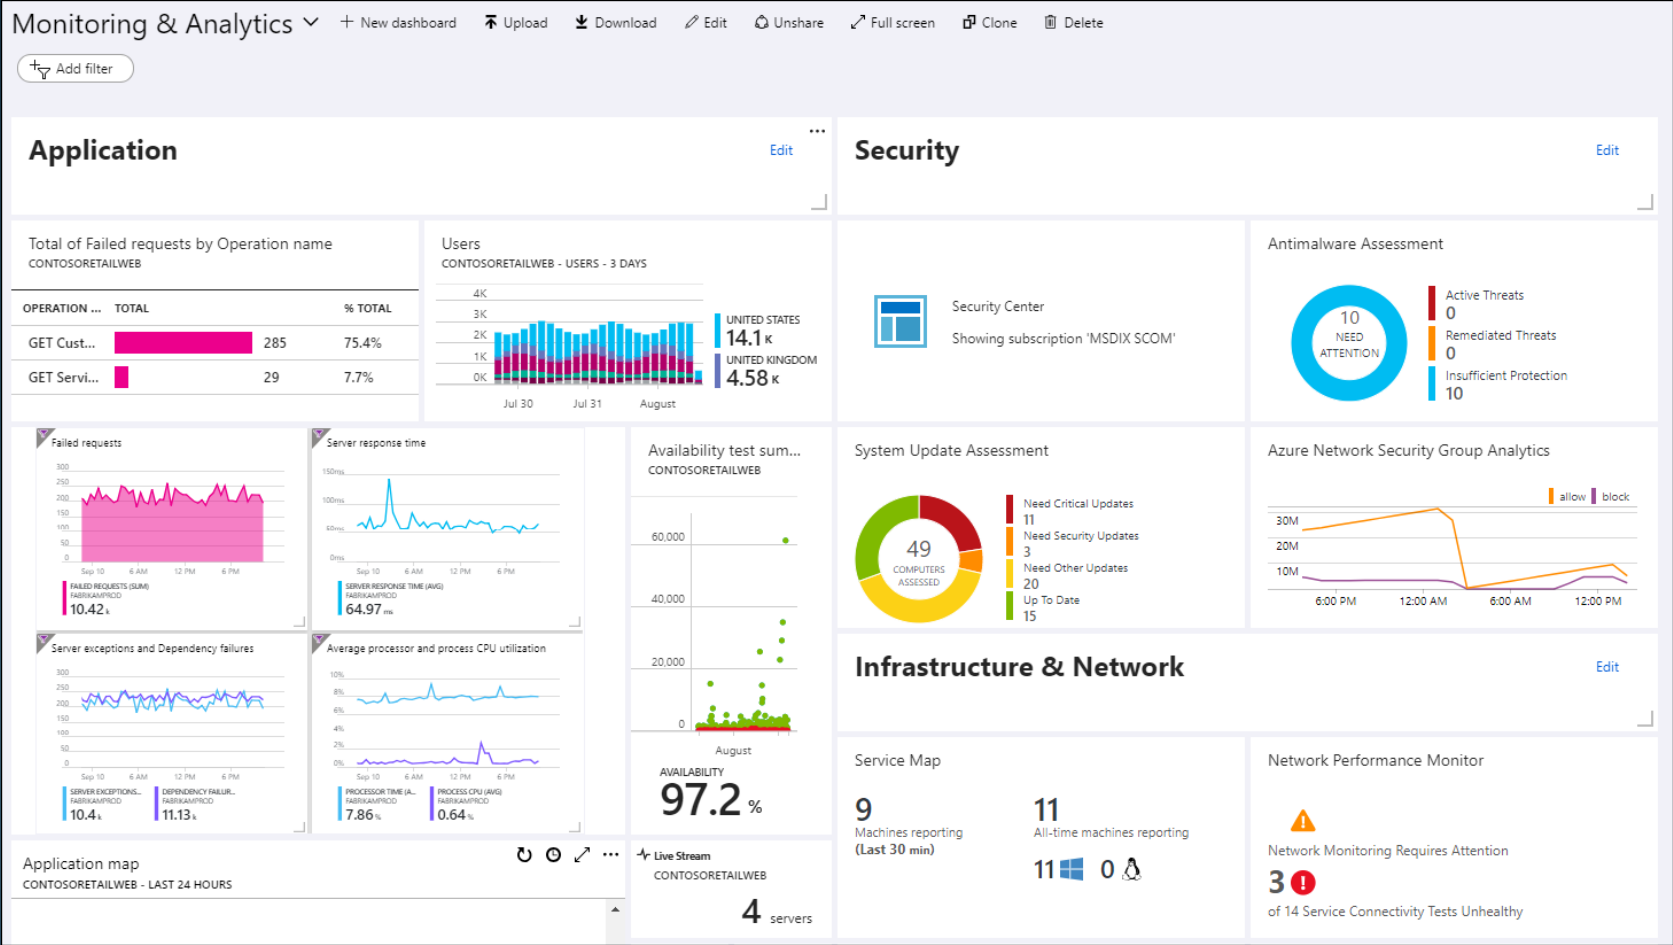
```
.create-or-alter function GetSysLogs(TimeWindow:string , Bucket:string )  
{  
cluster('help').database('SampleLogs').RawSysLogs  
| where timestamp > ago(totimespan(TimeWindow))  
| summarize LogCount=count() by name, bin(timestamp,  
totimespan(Bucket))  
| order by timestamp asc  
}
```

```
GetSysLogs('5d','1h')
```

# Analysis and reporting



## Dashboards in RTA - planned - to come...



# Thank you

Connect with me at:

 <https://linkedin.com/in/brianbonk>  
 <https://brianbonk.dk>  
 <https://github.com/brianbonk>



Session feedback





